

# Pivotal

A NEW PLATFORM FOR A NEW ERA

# Data Definition Language (DDL) in GPDB



Pivotal® Greenplum  
Database


# Agenda

- Introduction
- Database
- Schema, table
- Constraints
- Data types
- View, sequence, index, tablespace, trigger

# Databases – Overview

Greenplum Database instance:

- Supports multiple databases
- Creates three databases by default:
  - `template0`
  - `template1`
  - `postgres`



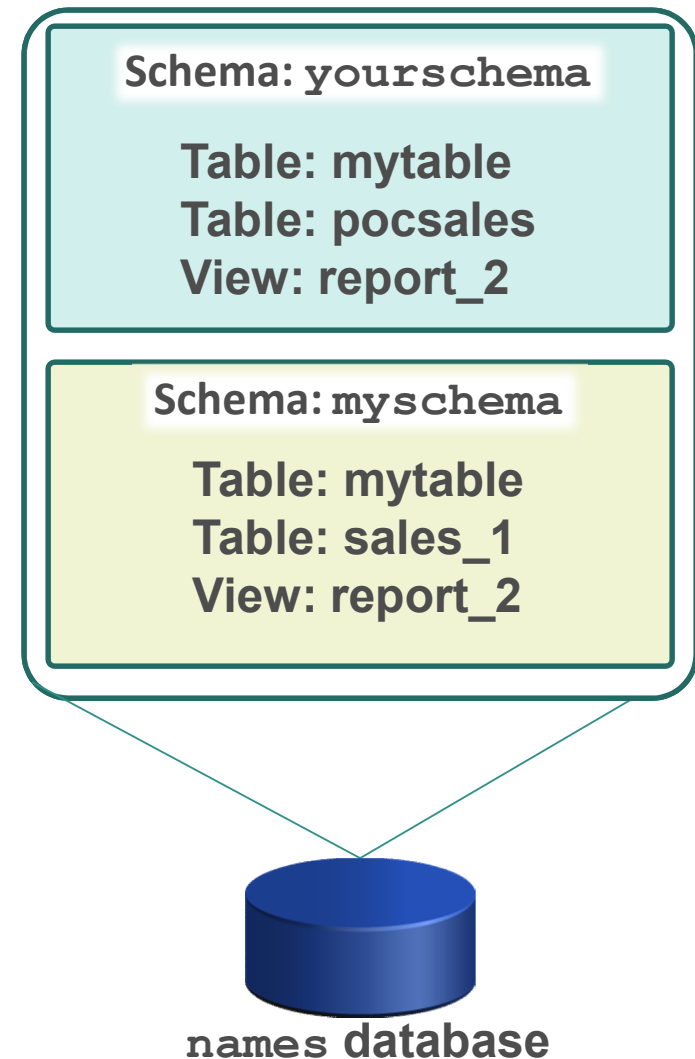
# Database SQL Commands

To manage databases, use the following SQL syntax or Greenplum application:

Action	SQL Syntax	Greenplum Application
Create a database	CREATE DATABASE	createdb
Drop a database	DROP DATABASE	dropdb
Alter a database: <ul style="list-style-type: none"><li>• Rename the database</li><li>• Assign a new owner</li><li>• Set configuration parameters</li></ul>	ALTER DATABASE	

# Schema – Overview

- Logically organize objects
  - A namespace
  - Do not represent users
  - Contain data objects
  - Use *qualified* names to access objects
- EXAMPLE:
- ```
myschema.mytable
```
- Can be added to the search path  
(search\_path)



# Pre-existing Schemas

The following schemas exist in every database:

- **PUBLIC** schema
- System level schemas:

| Schema                          | Description                                           |
|---------------------------------|-------------------------------------------------------|
| <code>pg_catalog</code>         | Stores system catalog names, functions, and operators |
| <code>information_schema</code> | Contains views that describe objects in the database  |
| <code>pg_toast</code>           | Stores large objects that exceed page size            |
| <code>pg_bitmapindex</code>     | Stores bitmap index objects                           |
| <code>pg_aoseg</code>           | Stores append-only objects                            |
| <code>gp_toolkit</code>         | Administrative schema                                 |

# Tables – Overview

## Tables in relational databases:

- Consist of columns and rows
- Have a fixed number and order of named columns
- Have a variable number of rows reflecting individual records
- No guaranteed row order

## Tables share the following characteristics:

- All tables in Greenplum Database are distributed based on a distribution key
- Some table features are not yet supported in Greenplum Database
  - Foreign key constraints
  - Limit on unique constraints



# Table Distribution Keys – Overview

- One column, or multiple columns, used to divide the data among the segments
- Should have high cardinality
- Specified with `CREATE TABLE` or `ALTER TABLE` using the `DISTRIBUTED BY` clause

```
CREATE TABLE sales
  (dt date, prc float, qty int, cust_id int,
   prod_id int, vend_id int)
DISTRIBUTED BY (dt, cust_id, prod_id);
```

- If table lacks a good candidate for distribution key, you can use `DISTRIBUTED RANDOMLY`
- If not explicitly declared, defaults to the table's `PRIMARY KEY` or the first column of the table

# Table SQL Commands – Modifying a Table

- Renaming columns
- Renaming tables
- Adding and removing columns

```
ALTER TABLE product ADD COLUMN msg_body TEXT;
```

- Adding and removing constraints

```
ALTER TABLE product ALTER COLUMN prod_no SET NOT NULL;
```

- Changing default values

```
ALTER TABLE product ALTER COLUMN prod_no SET DEFAULT -  
999;
```

- Changing column data types

```
ALTER TABLE products ALTER COLUMN price TYPE  
NUMERIC(10,2);
```

# Table Column Basics

- Each column has a data type
- Large set of data types
- User defined types supported
- Limit on the number of columns in a table: 1600

# Table and Column Constraints – Overview

Table and column constraints are supported with some limitations:

- CHECK **table** or column constraints
- NOT NULL column constraints
- UNIQUE column constraints
- PRIMARY KEY is used as the distribution key by default
- FOREIGN KEY constraints are supported, but *referential integrity* is **not** enforced

# Table and Column Constraints



## Example: Table and Column Constraints

```
CREATE TABLE TR_CONSTRAINTS (  
  transactionid int NOT NULL UNIQUE,  
  price numeric CHECK (price > 0),  
  on_sale char DEFAULT 'n'  
);
```

Column cannot  
be NULL

If no value is entered,  
a default value of 'n' is  
entered for that field

Column cannot  
have a value equal  
to or less than zero

Data in the column  
for that table is  
unique

# External Table SQL Commands

To manage external tables, use the following SQL syntax:

| Action                           | SQL Syntax                                                                                                                                                           |
|----------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Create a readable external table | <pre>CREATE EXTERNAL [WEB] TABLE<br/>LOCATION (...)   EXECUTE \'...\'<br/>FORMAT \'...\' (DELIMITER, \'...');</pre>                                                  |
| Create a writable external table | <pre>CREATE WRITABLE EXTERNAL [WEB]<br/>TABLE<br/>LOCATION (...)   EXECUTE \'...\'<br/>FORMAT \'...\' (DELIMITER, \'...')<br/>DISTRIBUTED BY (...)   RANDOMLY;</pre> |
| Drop an external table           | <pre>DROP EXTERNAL [WEB] TABLE</pre>                                                                                                                                 |

# Data Types – Overview

Data types:

- Constrain a column to storing only a specific kind of data
- In Greenplum are the same as data types in PostgreSQL
- Supports types defined in the SQL standard
- You can also define arrays, single- or multi-dimensional
- Can be created with the `CREATE TYPE SQL` command

For Greenplum distribution key columns:

- No geometric data types are supported
- No user-defined data types are supported

# Greenplum Database Data Types

| Name                    | Size                        | Description                                                                        |
|-------------------------|-----------------------------|------------------------------------------------------------------------------------|
| bigint                  | 8 bytes                     | Large range integer                                                                |
| bigserial               | 8 bytes                     | Large autoincrementing integer                                                     |
| bit [(n)]               | <i>n</i> bits               | Fixed-length bit string                                                            |
| bit varying [(n)]       | Actual # of bits            | Variable-length bit string                                                         |
| boolean                 | 1 byte                      | Logical boolean (true/false)                                                       |
| box                     | 32 bytes                    | Rectangular box in the plane<br>( <i>not allowed in distribution key columns</i> ) |
| bytea                   | 1 byte + binary             | Variable-length binary string                                                      |
| character [(n)]         | 1 byte + <i>n</i>           | Fixed-length, blank padded                                                         |
| character varying [(n)] | 1 byte + <i>string size</i> | Variable-length with limit                                                         |
| cidr                    | 12 or 24 bytes              | IPV4 and IPV6 networks                                                             |
| circle                  | 24 bytes                    | Circle in the plane<br>( <i>not allowed in distribution key columns</i> )          |



# Greenplum Database Data Types (Cont'd.)

| Name             | Size                 | Description                                                                        |
|------------------|----------------------|------------------------------------------------------------------------------------|
| date             | 4 bytes              | Calendar date (year, month, day)                                                   |
| decimal [(p,s)]  | variable             | User-specified precision, exact                                                    |
| double precision | 8 bytes              | Variable-precision, inexact                                                        |
| inet             | 12 or 24 bytes       | IPV4 and IPv6 hosts and networks                                                   |
| integer          | 4 bytes              | Usual choice for integer                                                           |
| interval [(p)]   | 12 bytes             | Time span                                                                          |
| lseg             | 32 bytes             | Line segment in the plane<br>( <i>not allowed in distribution key columns</i> )    |
| macaddr          | 6 bytes              | MAC addresses                                                                      |
| money            | 4 bytes              | Currency amount                                                                    |
| path             | 16+16 <i>n</i> bytes | Geometric path in the plane<br>( <i>not allowed in distribution key columns</i> )  |
| point            | 16 bytes             | Geometric point in the plane<br>( <i>not allowed in distribution key columns</i> ) |

# Greenplum Database Data Types (Cont'd.)

| Name                               | Size                        | Description                                                                       |
|------------------------------------|-----------------------------|-----------------------------------------------------------------------------------|
| polygon                            | 40+16 <i>n</i> bytes        | Geometric path in the plane<br>( <i>not allowed in distribution key columns</i> ) |
| real                               | 4 bytes                     | Variable-precision, inexact                                                       |
| serial                             | 4 bytes                     | Autoincrementing integer                                                          |
| smallint                           | 2 bytes                     | Small range integer                                                               |
| text                               | 1 byte + <i>string size</i> | Variable unlimited length                                                         |
| time [(p)] [without time zone]     | 8 bytes                     | Time of day only                                                                  |
| time [(p)] [with time zone]        | 12 bytes                    | Time of day only, with time zone                                                  |
| timestamp [(p)] [without timezone] | 8 bytes                     | Both date and time                                                                |
| timestamp [(p)] [with timezone]    | 8 bytes                     | Both date and time, with time zone                                                |

# Casting Data Types

A type cast:

- Is used to convert one data type to another
- Can be accomplished with the following syntax:

| Syntax                               | Example                                |
|--------------------------------------|----------------------------------------|
| <code>type 'string'</code>           | <code>REAL '2.117902'</code>           |
| <code>'string'::type</code>          | <code>'some text'::TEXT</code>         |
| <code>CAST ('string' AS type)</code> | <code>CAST ('2.117902' AS REAL)</code> |
| <code>value::type</code>             | <code>(1.0 / 3)::NUMERIC(3, 2)</code>  |

- Can be omitted if there is no ambiguity
- Accepts regular SQL notation or dollar quoting for string constants

# Additional Database Objects

Let us examine:

- Views
- Indexes
- Sequences
- Triggers
- Tablespaces

# Views – Overview

## Views:

- Let you save frequently used or complex queries
- Can be accessed with `SELECT` statement
- Are not materialized on disk
- Are managed and accessed with the following commands:

| Action                | Command                     |
|-----------------------|-----------------------------|
| Create a view         | <code>CREATE VIEW</code>    |
| Drop or remove a view | <code>DROP VIEW</code>      |
| List all views        | <code>\dv</code>            |
| See view definition   | <code>\dv+ view_name</code> |

# View Example

```
CREATE VIEW iot_json AS
SELECT
  (CASE WHEN a[1] = '' THEN '0.0' ELSE a[1] END)::FLOAT absolute_throttle_pos_b
, a[2]::FLOAT acceleration
, (CASE WHEN a[3] = '' THEN '0.0' ELSE a[3] END)::FLOAT accelerator_throttle_pos_d
, (CASE WHEN a[4] = '' THEN '0.0' ELSE a[4] END)::FLOAT accelerator_throttle_pos_e
, a[5]::FLOAT barometric_pressure
, a[7]::FLOAT catalyst_temp
, (CASE WHEN a[8] = '' THEN '0.0' ELSE a[8] END)::FLOAT control_module_voltage
, a[9]::FLOAT coolant_temp
, a[10]::FLOAT distance_with_mil_on
, TO_TIMESTAMP(SUBSTRING(a[27] FROM 1 FOR 10)::INT4)::TIMESTAMP date_time
FROM
(
  SELECT REGEXP_SPLIT_TO_ARRAY(csv_from_json(line), ',') a FROM iot_json_ext
) AS iot_row;
```

# Index – Overview

## Indexes:

- Use random seek to find a rows in a relational database
- Should be used sparingly in Greenplum Database
- Only used for highly selective queries
- Are not always favored at query runtime
- Should be checked to ensure they improve performance
- Should be dropped if not used in queries
- Are created automatically if using `PRIMARY KEY`

# Sequences – Overview

A sequence:

- Is used to auto-increment unique ID columns
- Is generated by the Greenplum Master sequence generator process (seqserver)
- Has some function limitations:
  - `lastval` and `currval` not supported
  - `setval` cannot be used in queries that update data
  - `nextval` sometimes grabs a block of values for some queries. So values may sometimes be skipped in the sequence if all of the block turns out not to be needed at the segment level.
  - Sequences are based on bigint arithmetic



**Note:** bigserial and serial data types are auto incrementing integers.



# Triggers – Overview

- Triggers are considered to be a type of function
- They automatically execute a particular function when a specific type of operation is performed (like update, delete, and insert)
- Have limited support in Greenplum, in that:
  - The function must be `IMMUTABLE` to execute on segment instances
    - The immutable function cannot execute any SQL or modify the database.
  - Triggers are not supported on Append-Optimized tables.

# Tablespaces

- Provide an alternate location for tables, partitions, temp tables, or indexes
- Can facilitate placing certain data on faster disk drives, for example
- Require superuser privileges to create or move data objects to them
- Are created using the following syntax:  

```
CREATE TABLESPACE tablespace_name  
FILESPACE filespace_name
```

# Review

- Database
- Schema, table
- Constraints
- Data types
- View, sequence, index, tablespace, trigger
- Work the lab!

# Pivotal

A NEW PLATFORM FOR A NEW ERA