

Greenplum Resource Management



Pivotal® **Greenplum**
Database

Greenplum Workload Management

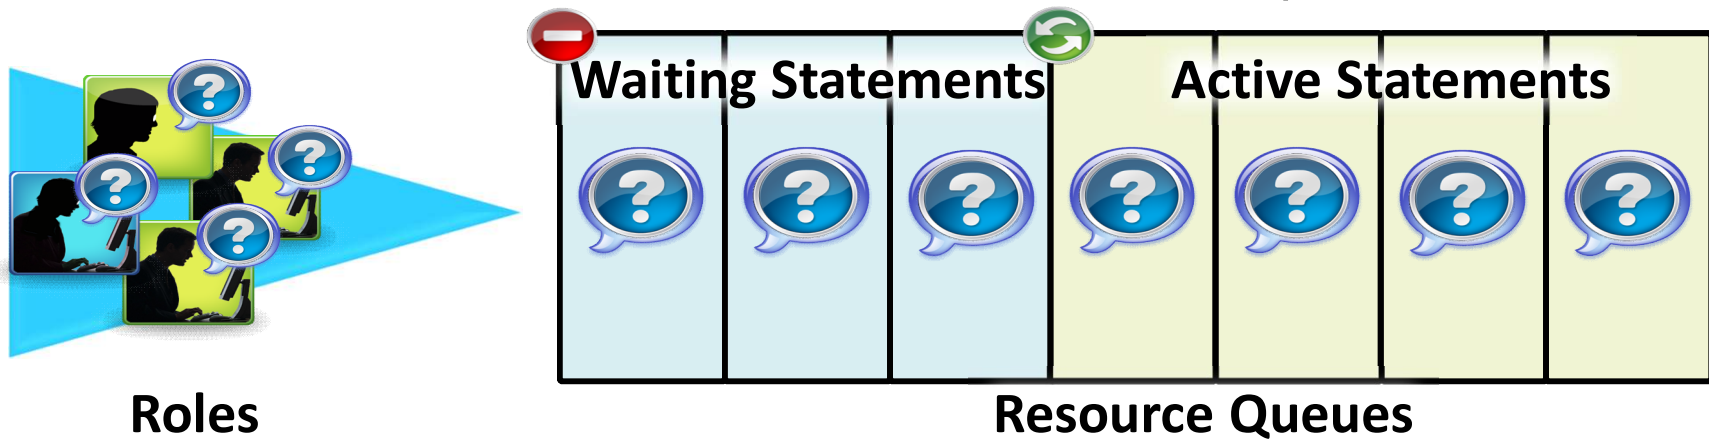
Let us examine the following:

- What is workload management?
- How do you create resource queues?
- How do you assign roles to resource queues?
- What is the runtime evaluation of resource queues?
- What are resource queue configuration parameters?
- How do you view the status of resource queues?

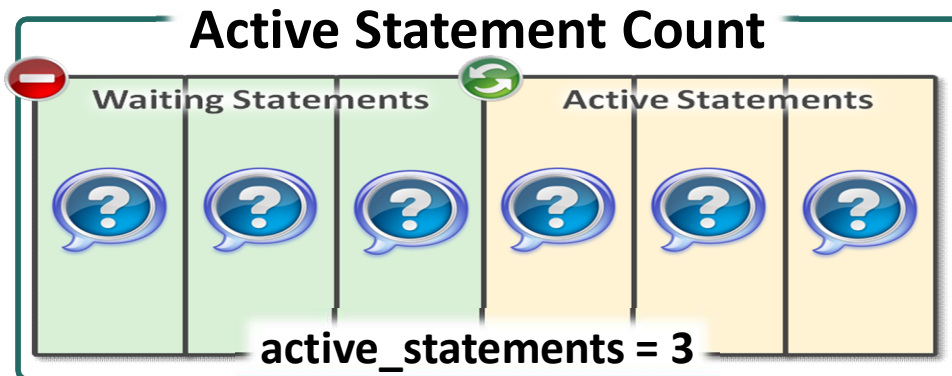
Workload Management Overview

Workload management:

- Is used to limit the number of active queries
- Is meant to prevent overloading of system resources
- Looks at resource allocation from statement-level point-of-view

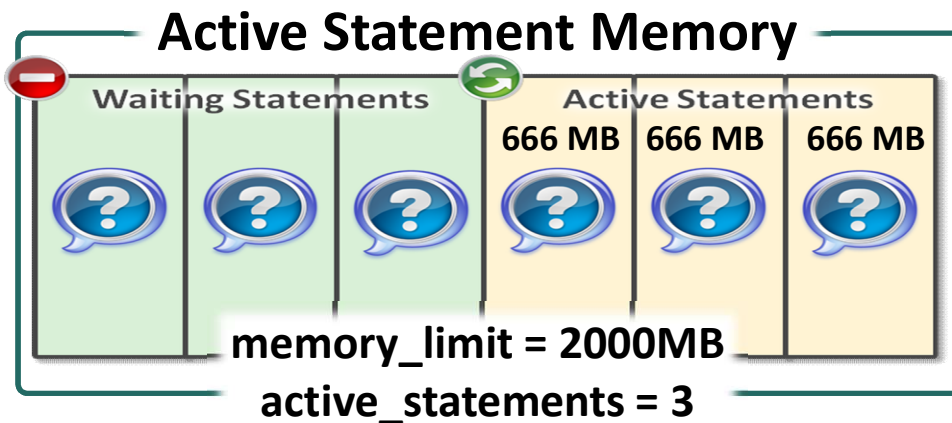


Configurable Limits on a Queue



This parameter determines query entry into the system to actively run.


The maximum number of queries executing in the active queue is limited by the `active_statements` limit.



This parameter determines resources allocated to the query.

Each query submitted to the queue will be given a portion of the memory. This amount is based on the amount of queries that can run in the queue or the cost of the query.

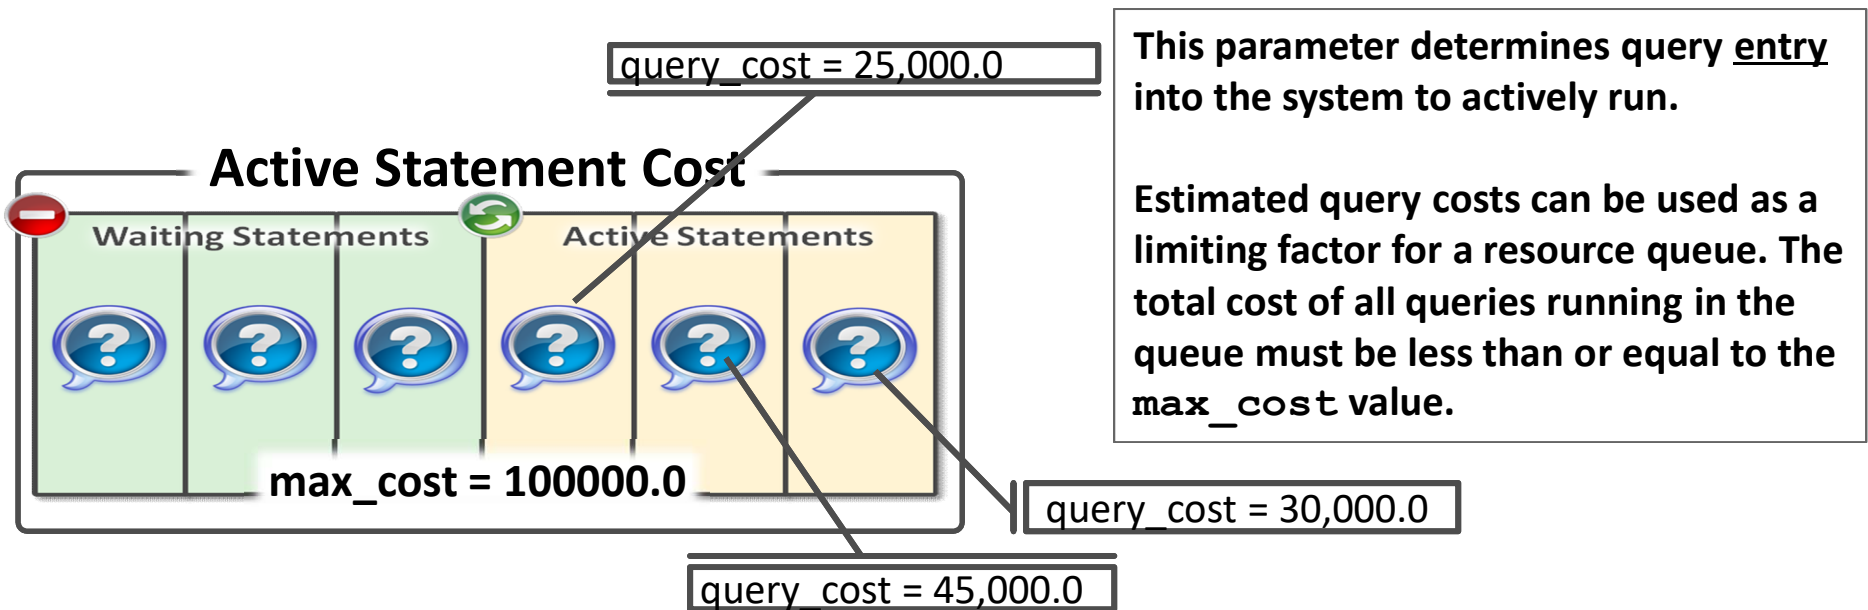
Configurable Limits on a Queue (Cont)



This parameter determines resources allocated to the query.

Queries with higher priorities consume more CPU resources over queries with lower priorities. Queries with the same priorities will evenly share CPU resources.

Configurable Limits on a Queue (Cont)



Creating Resource Queues

To create resource queues and manage thresholds, use the following commands:

Action	SQL Syntax
Create or alter a resource queue	CREATE ALTER RESOURCE QUEUE name WITH (MAX_COST=n ACTIVE_STATEMENTS =n [,COST_OVERCOMMIT] [,MIN_COST=n] [, MEMORY_LIMIT ='nM GB'] [, PRIORITY =level])
Drop a resource queue	DROP RESOURCE QUEUE name

Managing CPU Utilization with Priorities

Priorities:

- Control a resource queue's consumption of resources
- Are managed with the following SQL commands:

Action	SQL Syntax
Create a resource with a specific priority	<code>CREATE RESOURCE QUEUE <i>name</i> WITH (PRIORITY=<i>level</i>)</code>
Alter a resource queue to have a specific priority	<code>ALTER RESOURCE QUEUE <i>name</i> WITH (PRIORITY=<i>level</i>)</code>

- Can be one of the following:
 - MIN
 - LOW
 - MEDIUM (default setting if not specified)
 - HIGH
 - MAX

Managing Memory Allocation with Memory Limits

Query memory allocation is controlled with:

- MEMORY_LIMIT parameter for the resource queue:

Action	SQL Syntax	Memory Allocation per Query
Create or alter a resource with a specific memory limit and maximum number of statements	CREATE ALTER RESOURCE QUEUE <i>name</i> WITH (ACTIVE_STATEMENTS = <i>x</i> , MEMORY_LIMIT='xxxM GB')	MEMORY_LIMIT / ACTIVE_STATEMENTS
Create or alter a resource queue to have a specific memory limit with MAX_COST defined	CREATE ALTER RESOURCE QUEUE <i>name</i> WITH (MAX_COST = <i>x.x</i> , MEMORY_LIMIT='xxxM GB')	MEMORY_LIMIT * (query_cost / MAX_COST)

- At the database level with statement_mem and max_statement_mem

Creating Resource Queues – Examples



Example: Create a resource queue with up to 3 active queries

```
CREATE RESOURCE QUEUE adhoc with  
(ACTIVE_STATEMENTS=3);
```



Example: Create a resource queue with a planner cost limit

```
CREATE RESOURCE QUEUE webuser WITH  
(MAX _COST=100000.0);
```



Example: Create a resource queue with a minimum cost limit

```
CREATE RESOURCE QUEUE adhoc WITH  
(ACTIVE_STATEMENTS=10, MIN_COST=100.0);
```

Assigning Roles to Resource Queries

Resource queues:

- Must be assigned at the user-level (group-level ignored)
- Do not affect superuser roles

Action	SQL Syntax
Add a resource queue to a role	<code>ALTER ROLE role_name RESOURCE QUEUE queue_name;</code>
Create a role and assign it to a resource queue	<code>CREATE ROLE role_name WITH LOGIN RESOURCE QUEUE queue_name;</code>



Note: All users are assigned to a resource queue, even if one is not specified. The default resource queue is `pg_default`.

Runtime Evaluation of Resource Queues

For runtime evaluation:

- Resource queues are evaluated independently of each other
- Superusers (and unassigned roles) are exempt
- Queries are evaluated on first-in, first-out basis

If a query causes the queue to exceed its limits:

- Query must wait until queue resources are free
- Query must wait until queue is idle (`COST_OVERCOMMIT=TRUE`)
- Query will never run (`COST_OVERCOMMIT=FALSE`)

Evaluated SQL statements include:

- `SELECT`, `SELECT INTO`, `CREATE TABLE` **AS** `SELECT`,
`DECLARE CURSOR`
- (optional) `INSERT`, `UPDATE`, `DELETE`

Resource Queue Configuration Parameters

The following parameters are used to configure resource queues:

Configuration Parameter	Default Value	Description
<code>max_resource_queues</code>	9	The maximum number of resource queues in the system
<code>max_resource_portals_per_transaction</code>	64	The maximum number of open cursors per transaction
<code>resource_select_only</code>	on	Determines which queries are managed by resource queues
<code>stats_queue_level</code>	off	Enables the collection of statistics on resource queue usage
<code>resource_cleanup_gangs_on_wait</code>	on	Clean up idle segment worker processes

Memory Utilization System Parameters

Configuration Parameter	Default Value	Description
<code>gp_resqueue_memory_policy</code>	<code>eager_free</code>	The query plan is divided into stages and Greenplum Database eagerly frees memory allocated to a previous stage at the end of that stage's execution, then allocates the eagerly freed memory to the new stage
<code>statement_mem</code> <code>max_statement_mem</code>	125 (MB) 2000 (MB)	<code>statement_mem</code> : Allocates segment host memory per query <code>max_statement_mem</code> : Sets the maximum memory limit for a query. $(\text{seg_host_physical_memory}) / (\text{average_number_concurrent_queries})$
<code>gp_vmem_protect_limit</code>	8192 (MB)	Upper memory boundary that all query processes can consume on a segment host
<code>gp_vmem_idle_resource_timeout</code>	18000 (Milliseconds)	If database session is idle for longer than the time specified, the session will free system resources (such as shared memory), but remain connected to the database
<code>gp_vmem_protect_segworker_cache_limit</code>	500 (MB)	If a query executor process consumes more than this configured amount, then the process will not be cached for use in subsequent queries after the process completes

Viewing Status of Resource Queues

The following query system tables provide insight into resource queues:

System Table	Description
pg_resqueue	Resource queues and attributes
gp_toolkit.gp_resq_role pg_roles	Role to resource queue assignments
pg_locks	Queues that have waiting statements
pg_stat_activity	Process information about active and waiting queries

To view queue status and statistics, access the following tables:

System Table	Description
gp_resqueue_status pg_resqueue_status	Queue limits, number of active and waiting queries
pg_stat_resqueues	Queue statistics and performance over time

Wrapping Up

- Roles and privileges
- Role assignment to object privileges
- Security issues that can affect your database and corresponding data
- Creating roles with specific privileges to control access
- Isolating users at the physical and functional layers
- Workload management process