Greenplum Workload Manager

Pivotal® Greenplum Database

Pivotal

Agenda

- Workload Manager Overview
- Workload Management
  - Services and Agents
  - Actions, Conditions, and Rule Framework

Learning Objectives

- Learn how to Configure Limits
- Enable Workload Management

Workload Manager (WLM)

Workload Manager Motivation

- Enables operators to maximize value of investment
- When the system is under load, WLM can throttle processes or kill low-priority queries based on operator-defined rules
- Conserves system resources for high-priority operations

Workload Manager

- Tool to monitor and manage queries
  - Monitor Greenplum Database queries and host utilization statistics
  - Log when a query exceeds a threshold
  - Throttle the CPU usage of a query when it exceeds a threshold
  - Terminate or cancel a query
  - Detect memory, CPU, or disk I/O skew occurring during the execution of a query
  - Create detailed rules to manage queries

  Workload Manager Docs

Pivotal

Workload Manager Installation

Run the gp-wlm.bin installer on the GPDB Master Server

- Installs on the GPDB Master Server
- Automatically distributes software to all Segment Servers
- When upgrading, automatically rolls back if upgrade fails

```
./gp-wlm.bin --install=<DIR> [ --force ] [ --install-concurrency=<COUNT> ]
[ --no-remove-old ] [ --skip-health-check ] [ --dbname-records=<database_name> ]
[ --tool-manifest=<FILE> ]
```

- Installation Sanity Check with "cluster-health-check" utility (automatically run by installer)

Pivotal

Services

- Installs and runs five services
  - agent
  - cfgmon
  - rabbitmq
  - rulesengine
  - svcmon (service monitor)
- Command line **interactive** mode with "gp-wlm"
- Graphical interface (curses interface) "gptop”

Manage Services

- Manage with "svc-mgr.sh" where

```
INSTALLDIR/gp-wlm/bin/svc-mgr.sh \
 --service=SVCNAME \
 --action=ACTION
SVCNAME = agent | cfgmon | rabbitmq | rulesengine | all
ACTION (defaults to local unless prefixed with "cluster-" )
ACTION = start | cluster-start
stop | cluster-stop
status | cluster-status
restart | cluster-restart
enable |cluster-enable
disable |cluster-disable
```

Using Workload Manager Rules

- Trigger Actions on events
- Rule specifies action to execute when specific conditions are detected
- Best Practice guideline – watch your logs first to determine what rules make sense for you
  - Problem queries
  - Throttle queries that consume too much CPU
  - Terminate queries that disrupt DB operations

Actions

- Acts like a "do <action-exp> WHEN <condition-exp>"
- Action Expressions
  - `gpdb_record`: record a custom message and details of the database query process in the gp_wlm_records database table
  - `host:throttle_gpdb_query:` throttle a Greenplum Database query on a specified host
  - `host:pg_cancel_backend:` cancel the current query in a session on a host by calling the PostgreSQL pg_cancel_backend() function
  - `pg_terminate_backend:` terminate a session by calling the PostgreSQL pg_terminate_backend() function

Conditions

- Greenplum scoped-datum
  - E.g. session_id:host:pid:runtime
    scoping = session_id:host:pid
    datum = runtime
- Example of an Action / Condition

  pg_terminate_backend() when session_id:host:pid:runtime > 120

- Datums are data items collected by the agent, and include operating system statistics, OS process statistics, and database query data

List of All Datums, scopes, and data formats

Pivotal

Rule framework

- Add Rule
  `rule add [transient] <name> <action-name>(<action-args>) when <expression> [including <datum_list>]`
- Show Rules – `rule show { all | rule-name }`
- Delete Rule - `rule delete rule-name`
- Modify Rule - `rule modify [transient] name action-name (action-args) when expression`
- Save Rule to Disk - `rule dump path`
- Import Rules from Disk - `rule import path`
- Restore Rules from Disk - `rule restore path`

Example Rules

log_query gpdb_record(message="Query > 10 seconds") when session_id:host:pid:runtime > 10 and session_id:host:pid:current_query != '<IDLE>'

kill_query pg_terminate_backend() when session_id:host:pid:runtime > 1200 and session_id:host:pid:current_query != '<IDLE>'

throttle_q host:throttle_gpdb_query(max_cpu=10) when session_id:host:pid:current_query =~ /.*select count.*/

record_skew gpdb_record(message="Skew") when session_id:disk_write_bytes_per_sec_skew >= 90 and session_id:host:pid:usename =~/.*/ and (session_id:host:pid:current_query =~ /.*insert.*/ or session_id:host:pid:current_query =~ /.*create.*select./)

Pivotal

Wrapping Up

In this module we covered:

- An overview of Workload Manager
- The different services and agents that are installed
- The concept of Actions and Conditions
- And, how to set up rules