

Pivotal

A NEW PLATFORM FOR A NEW ERA

GPDB Storage Considerations

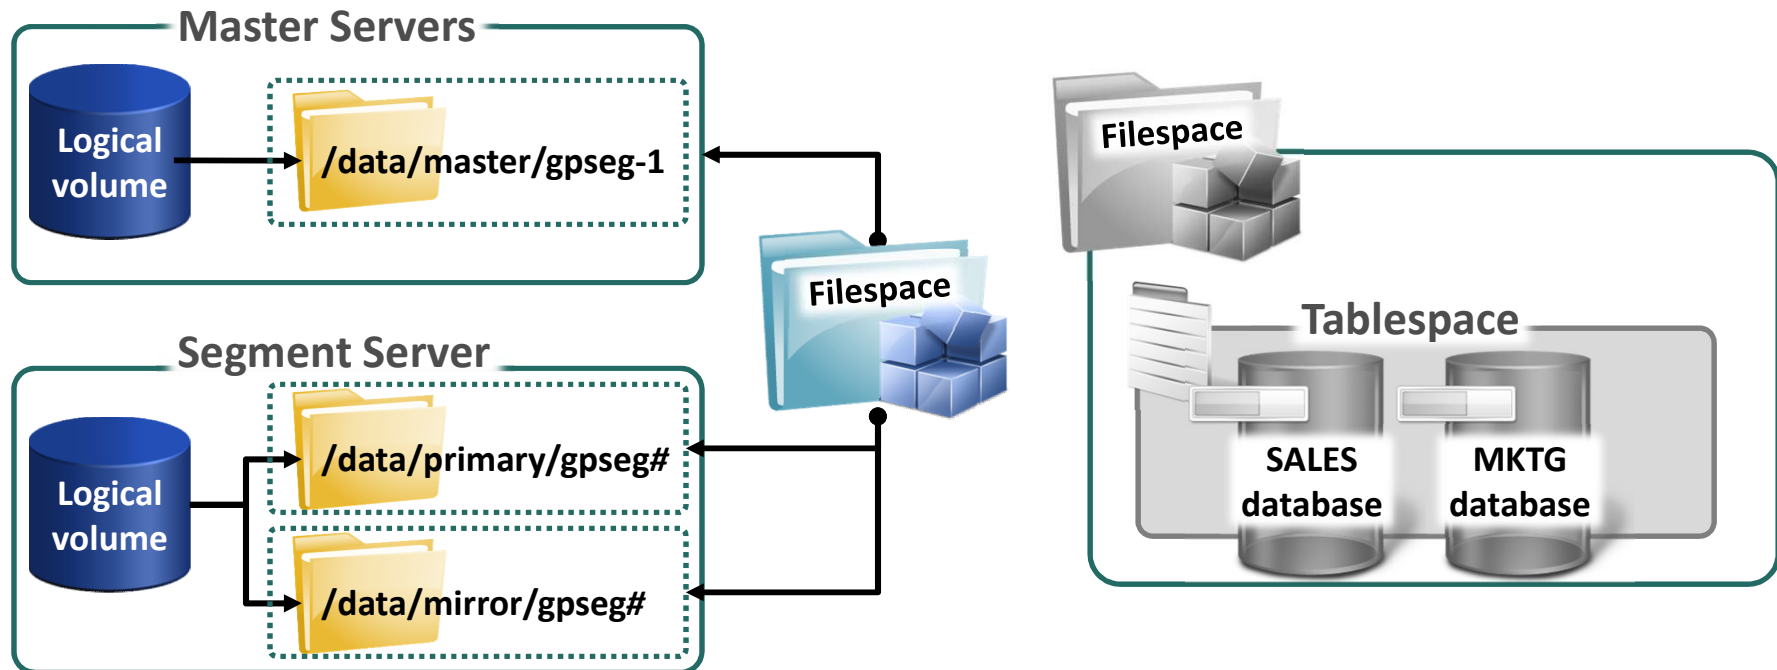


Pivotal® **Greenplum**
Database

Agenda

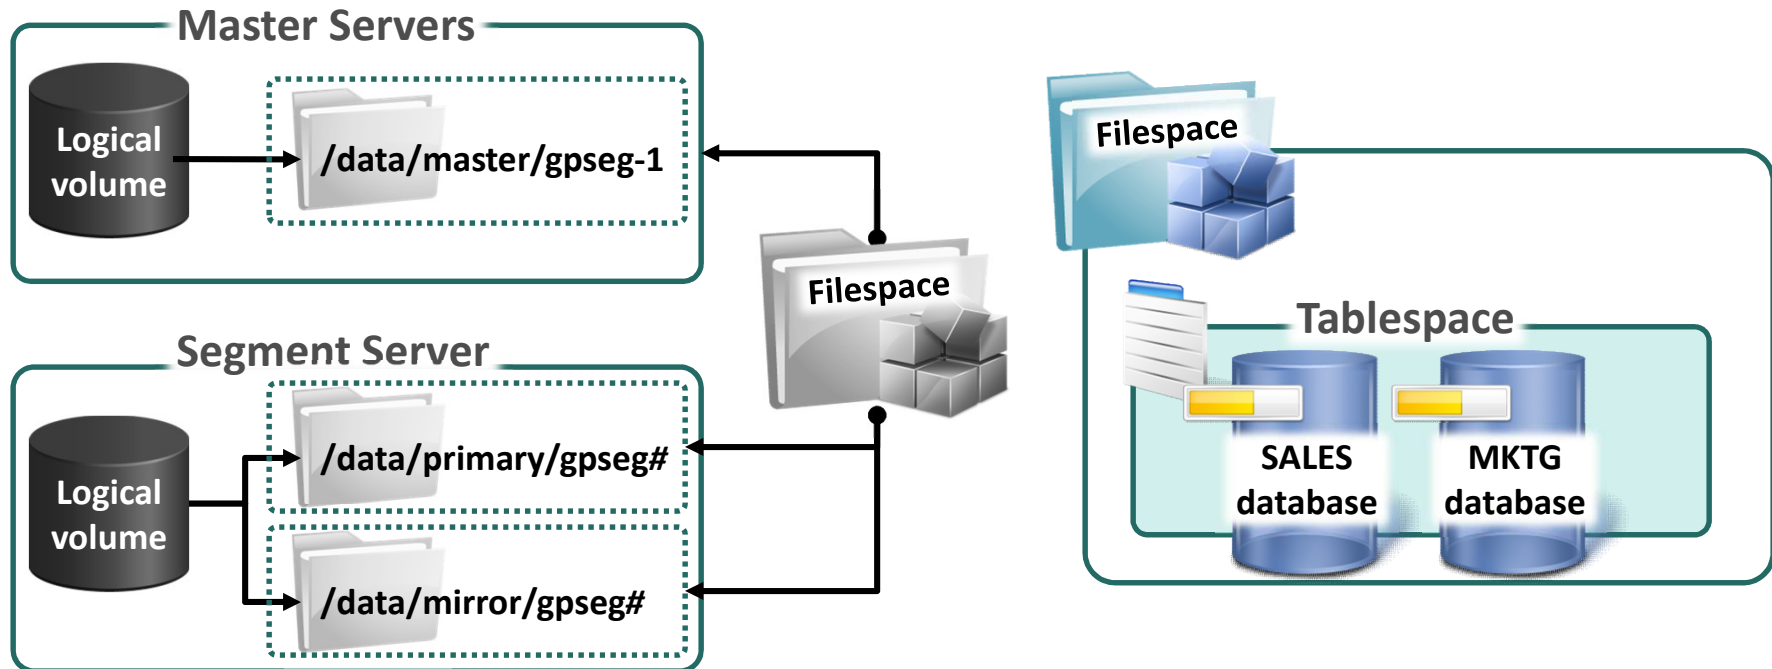
- Introduction
- Tablespaces and tablespaces
- Additional table types (external, temp)
- Table storage models
- Compression options
- Test it out in the lab

Filespaces



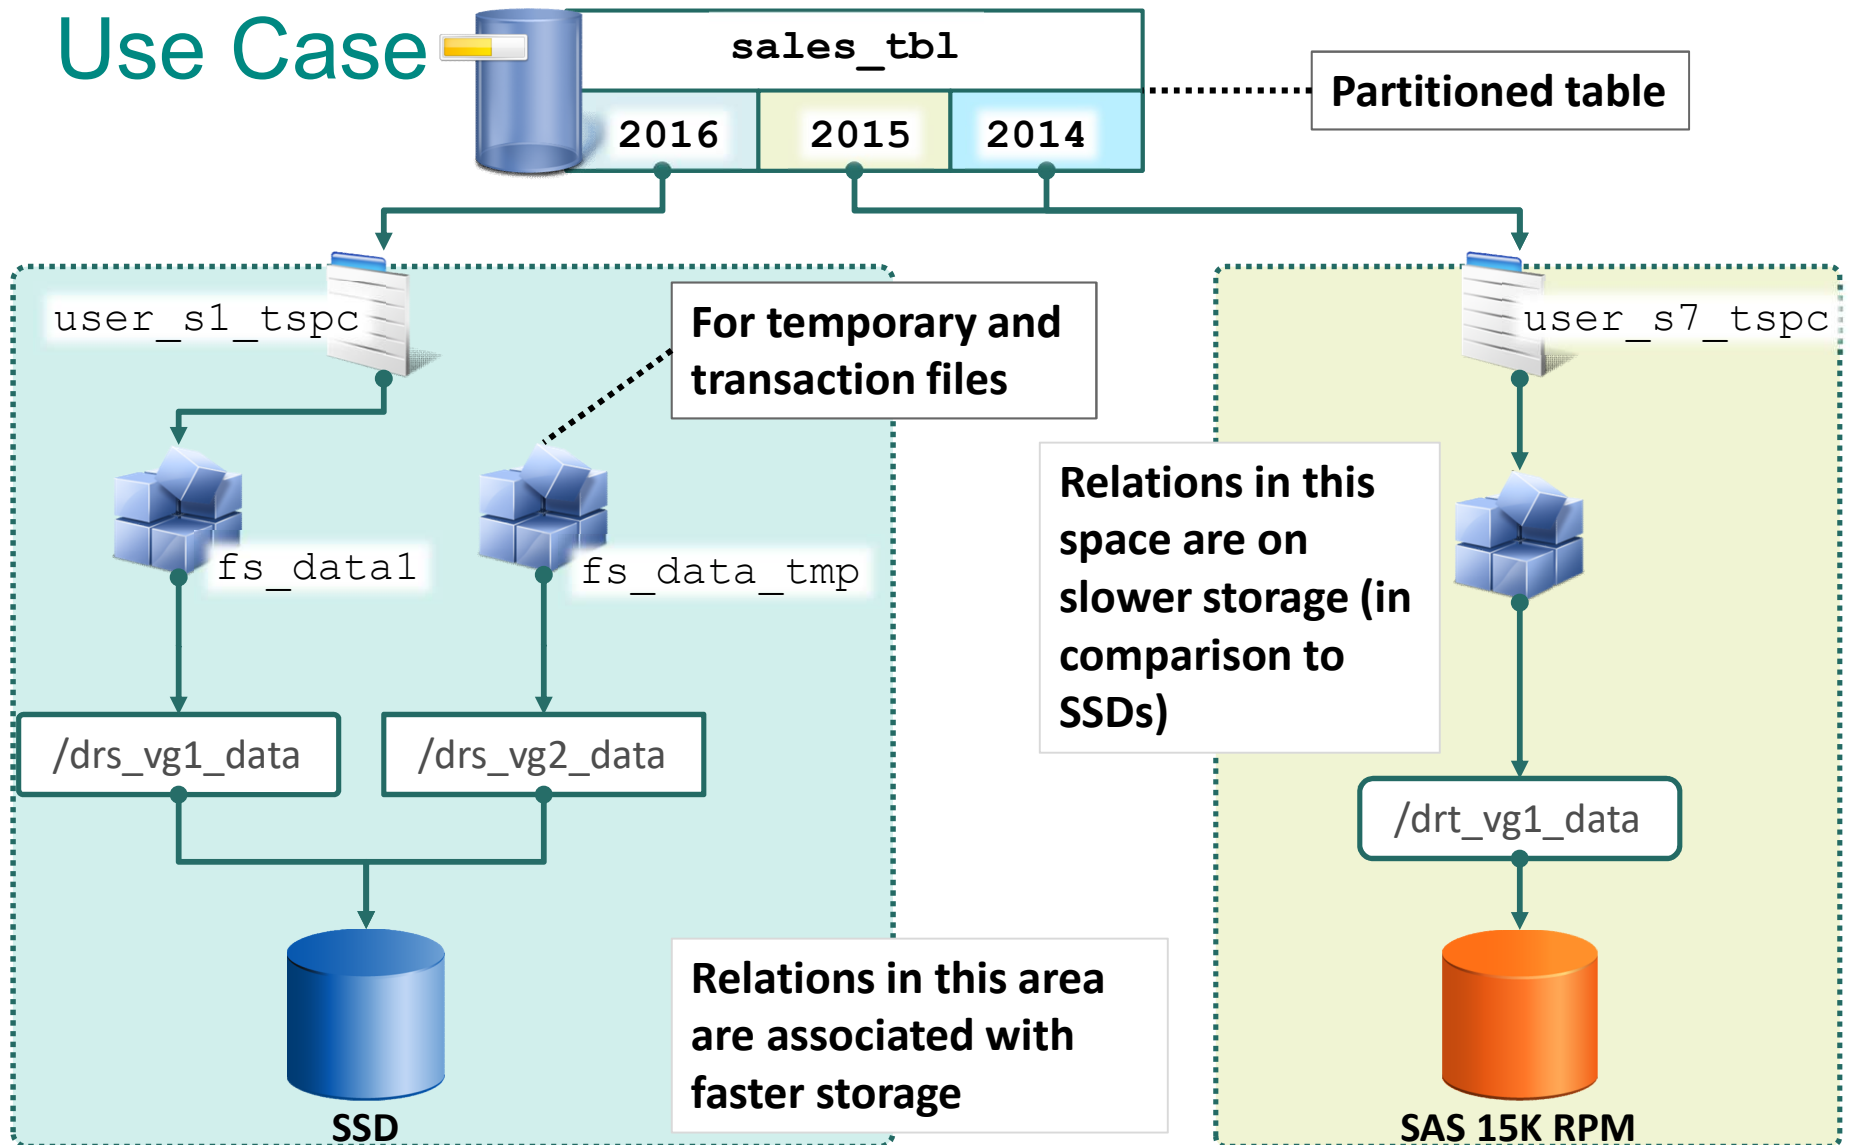
- By default, the system filesystem, `pg_system`, is created on initialization
- All system relations are stored in the system filesystem by default
- All user relations are also stored in the system filesystem by default

Tablespaces



- Tablespaces *sit atop* filespaces interacting with the underlying filesystem
- A filespace can support multiple tablespaces
- Two tablespaces are created on initialization: `pg_default` and `pg_global`

Filespace and Tablespace Implementation – Use Case



Creating the Filespace Configuration File

```
gpadmin@mdw:~$ gpfilespace -o gpfilespace_config
20131010:15:23:22:003198 gpfilespace:mdw:gpadmin-[INFO]:-
A tablespace requires a file system location to store its database
files. A filespace is a collection of file system locations for all components
in a Greenplum system (primary segment, mirror segment and master instances).
Once a filespace is created, it can be used by one or more tablespaces.

20131010:15:23:22:003198 gpfilespace:mdw:gpadmin-[INFO]:-getting config
Enter a name for this filespace
> fs_data1

Checking your configuration:
Your system has 2 hosts with 1 primary and 1 mirror segments per host.
Your system has 2 hosts with 0 primary and 0 mirror segments per host.

Configuring hosts: [sdw2, sdw1]

Please specify 1 locations for the primary segments, one per line:
primary location 1> /data/user_spc/primary

Please specify 1 locations for the mirror segments, one per line:
mirror location 1> /data/user_spc/mirror

Configuring hosts: [smdw, mdw]

Enter a file system location for the master
master location> /data/user_spc/master
20131010:15:23:51:003198 gpfilespace:mdw:gpadmin-[INFO]:-Creating configuration file.
20131010:15:23:51:003198 gpfilespace:mdw:gpadmin-[INFO]:-[created]
20131010:15:23:51:003198 gpfilespace:mdw:gpadmin-[INFO]:-
To add this filespace to the database please run the command:
gpfilespace --config /home/gpadmin/gpfilespace_config

[gpadmin@mdw ~]$
```

All directories must exist and be owned by gpadmin

Directories must already exist on segment hosts

Directory must already exist on master and standby master hosts

Creating the Filespace



Example: Filespace Configuration File

Configuration file contains all directories needed by masters and segments

```
$ cat gpfilespace_config
filespace:fs_data1
mdw:1:/data/user_spc/master/gpseg-1
smdw:6:/data/user_spc/master/gpseg-1
sdw2:3:/data/user_spc/primary/gpseg1
sdw2:4:/data/user_spc/mirror/gpseg0
sdw1:2:/data/user_spc/primary/gpseg0
sdw1:5:/data/user_spc/mirror/gpseg1
```

```
gpadmin@mdw:~$ gpfilespace --config /home/gpadmin/gpfilespace_config
20131010:15:52:18:003805 gpfilespace:mdw:gpadmin-[INFO]:-
A tablespace requires a file system location to store its database
files. A filespace is a collection of file system locations for all components
in a Greenplum system (primary segment, mirror segment and master instances).
Once a filespace is created, it can be used by one or more tablespaces.

20131010:15:52:18:003805 gpfilespace:mdw:gpadmin-[INFO]:-getting config
Reading Configuration file: '/home/gpadmin/gpfilespace_config'
20131010:15:52:18:003805 gpfilespace:mdw:gpadmin-[INFO]:-Performing validation on paths
.....

20131010:15:52:19:003805 gpfilespace:mdw:gpadmin-[INFO]:-Connecting to database
20131010:15:52:19:003805 gpfilespace:mdw:gpadmin-[INFO]:-Filespace "fs_data1" successfully
created
[gpadmin@mdw ~]$
```

Create the filespace as the gpadmin user using the filespace configuration file

Creating the Tablespace




Usage: Syntax to create a tablespace

```
gpadmin=# CREATE TABLESPACE tablespace_name [OWNER username]  
FILESPACE filespace_name;
```

Must be unique

```
gpadmin@mdw:~  
gpadmin=# select * from pg_filespace;  
  fsname  | fsowner  
-----+-----  
pg_system |      10  
fs_data1  |      10  
(2 rows)  
  
gpadmin=# create tablespace user_s1_tspc filespace fs_data1;  
CREATE TABLESPACE  
gpadmin=# select spcname, fsname from pg_tablespace,pg_filespace where pg_filespace.oid=pg_  
_tablespace.spcfsoid;  
  spcname  | fsname  
-----+-----  
pg_default | pg_system  
pg_global  | pg_system  
user_s1_tspc | fs_data1  
(3 rows)  
  
gpadmin=#
```

The tablespace has
now been successfully
created



Note: The maximum number of tablespaces and filespace are represented as `gp_max_tablespaces` and `gp_max_filespace` in the master `postgresql.conf` file.


Applying the Tablespace

Object	Example
Database	<code>CREATE DATABASE tt_db TABLESPACE user_s1_tspc;</code>
Table	<code>CREATE TABLE tt_rt (id int) TABLESPACE user_s1_tspc;</code>
Partitioned Table	<code>CREATE TABLE ttct2_part_rt (id int, id2 int) PARTITION BY LIST (id) (PARTITION one VALUES (1), PARTITION two VALUES (2) TABLESPACE user_s1_tspc, PARTITION three VALUES (3));</code>
Index	<code>CREATE INDEX tt_idx ON tt_rt (id) TABLESPACE user_s1_tspc;</code>



Note: The default tablespace for the environment can be set by modifying the `default_tablespace` parameter in the master server's `postgresql.conf` file.

Additional Table Types



Temporary tables can be used for:

- Storing transient results needed for other session queries
- Perform transformations on data

External tables:

- Facilitate parallel data loading
- Stream data in from external sources
- Push data out of the database, in parallel

Temporary Tables – Overview

- Session-specific
- Dropped at the end of the session
- Take precedence over permanent tables of the same name
- Created in a special schema created on connection to a session
- Are distributed
- Can be indexed and analyzed

Creating a Temporary Table



Example: Creating a temporary table

```
gpadmin=# CREATE TEMP[ORARY] TABLE monthlytranssummary (  
    storeid    INTEGER,  
    customerid INTEGER,  
    transmonth SMALLINT,  
    salesamttot DECIMAL(10,2)  
)  
ON COMMIT PRESERVE ROWS  
DISTRIBUTED BY (storeid, customerid);
```

You can define how the temporary table will be handled *for transactions* with the ON COMMIT clause

The following options to the ON COMMIT clause let you define how a temporary table is handled:

- PRESERVE ROWS – No action is taken on the table
- DELETE ROWS – The table is truncated
- DROP – The table is dropped

Temporary Tables – Two Use Cases

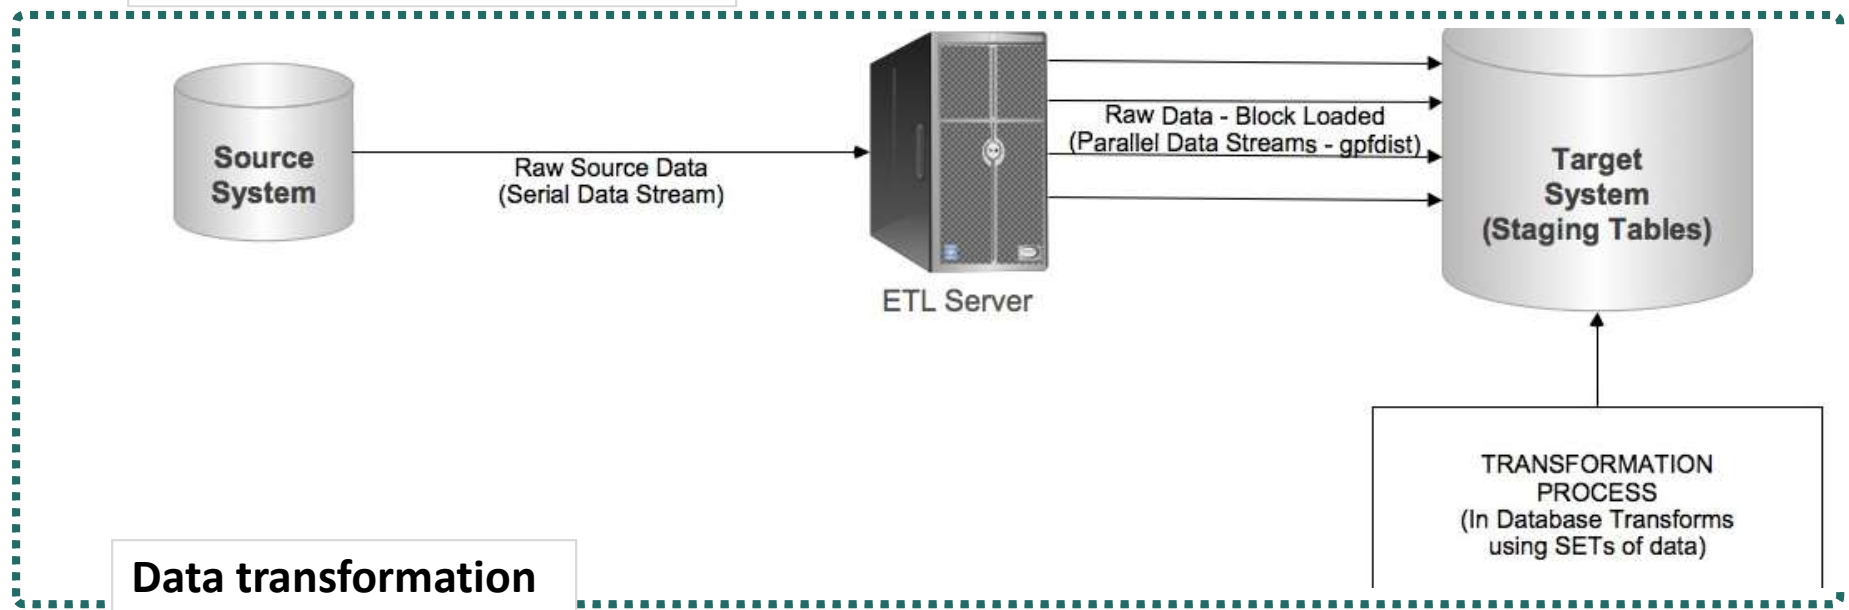
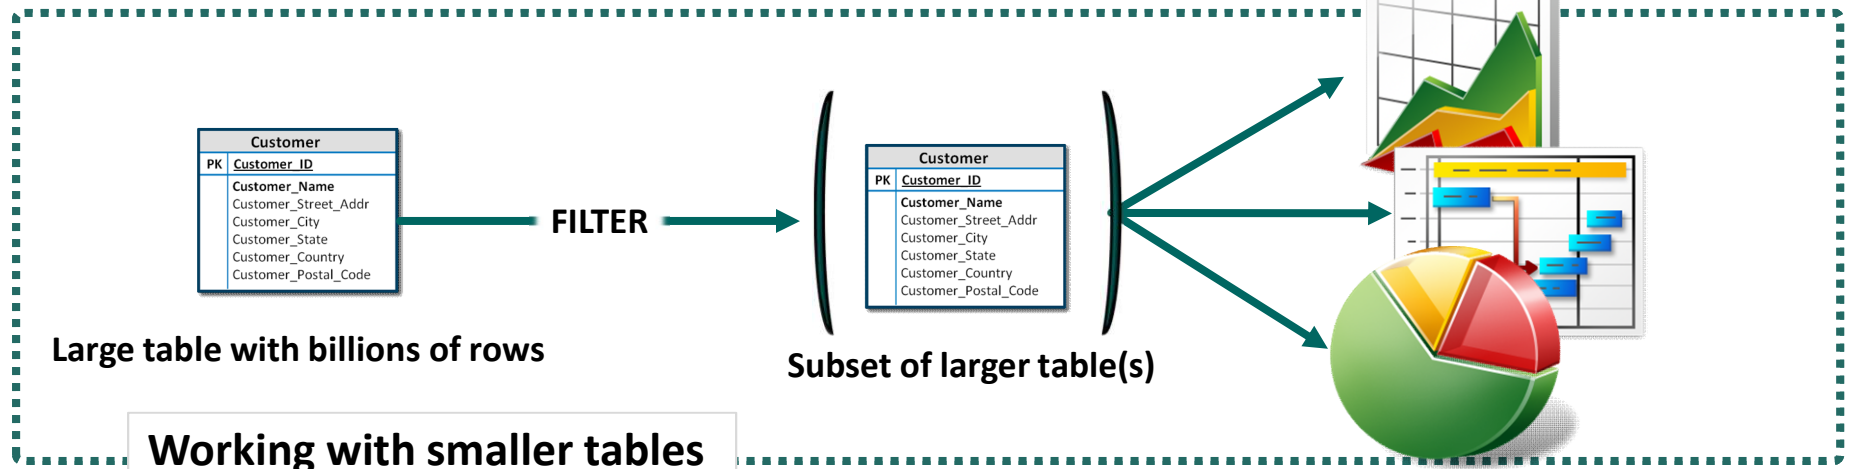



Table Storage Models



Customer
Customer_ID
Customer_Name
Customer_Street_Addr
Customer_City
Customer_State
Customer_Country
Customer_Postal_Code

Heap storage

- Default storage model
- Supports INSERT, UPDATE, DELETE
- Best for:
 - Data that is often modified
 - Smaller dimension tables
- Supports row-oriented tables
- Uses MVCC to support transactions




Customer
Customer_ID
Customer_Name
Customer_Street_Addr
Customer_City
Customer_State
Customer_Country
Customer_Postal_Code

Append-optimized storage

- Append-optimized storage model:
- Optimized for data warehouses
- Works best with denormalized data
- Supports UPDATE and DELETE
- Best for:
 - Older data
 - Large fact tables
- Supports row and column-oriented tables
- Supports in-database compression
- Uses a Visibility Map (visimap) to hide outdated rows


Row-Oriented and Column-Oriented Tables

Row-oriented storage



- Supports mixed workloads (INSERT, UPDATE, DELETE, SELECT)
- Is supported with on both heap and append-optimized storage

Column-oriented storage



- Works well with data warehouse workloads
- Works well for data where you aggregate over a small number of columns
- Efficient for data where you modify a single column
- Supported on append-optimized storage

Creating Heap and Append-Optimized Tables

Action	Example
Creating a heap, row-oriented table	<pre>CREATE TABLE tc_heap (id int, descr text) DISTRIBUTED BY (id);</pre>
Creating an append-optimized, row-oriented table	<pre>CREATE TABLE tc_ao (id int, sales float) WITH (appendonly=true) DISTRIBUTED BY (id);</pre>
Creating an append-optimized, column-oriented table	<pre>CREATE TABLE tc_ao_c (id int, sales float) WITH (appendonly=true, orientation=column) DISTRIBUTED BY (id);</pre>



Note: You cannot modify the storage or orientation of a table once defined. You can create a new table with the desired options and migrate your data.

Compressing Table Data

Compression Algorithm	Compression Levels	Description	Table-Level Compression	Row-Level Compression
ZLIB	1 - 9	Offers the most compact ratio with a potential impact to CPU performance	Supported	Supported
QUICKLZ	1	Offers faster, but lower, data compression	Supported	Supported
RLE_TYPE Delta Compression (specific data types)	1 - 4	Offers run-length encoding compression for columns based on repeated values	Unsupported	Supported

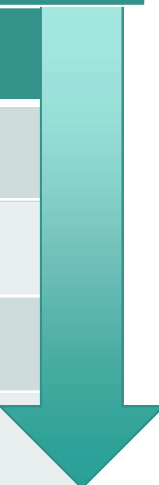


Question: What type of data do you think would work well with the different offerings of compression?

Defining Append-Optimized Compression Tables

Action	Example
Creating a zlib compressed table with compression level 5	<pre>CREATE TABLE tc_ao_zlib5 (id int, sales float) WITH (appendonly=true, compresstype=zlib, compresslevel=5) DISTRIBUTED BY (id);</pre>
Creating a quicklz compressed table	<pre>CREATE TABLE tc_ao_quicklz (id int, sales float) WITH (appendonly=true, compresstype=quicklz) DISTRIBUTED BY (id);</pre>
Creating an AO table with an RLE compressed column and a zlib compressed column	<pre>CREATE TABLE tc_ao_rletype (id int, sales float ENCODING (compresstype=zlib, compresslevel=3), salesdate date ENCODING (compresstype=rle_type)) WITH (appendonly=true, orientation=column) DISTRIBUTED BY (id);</pre>

Defining Default Table Storage Options

gp_default_storage_options			 Highest priority
Options	Level	Command	
APPENDONLY	Object level	CREATE TABLE ... WITH (...)	
BLOCKSIZE	Role level	ALTER ROLE ... SET ...	
CHECKSUM			
COMPRESSTYPE	Database level	ALTER DATABASE ... SET ...	
COMPRESSLEVEL	System level		
ORIENTATION		gpconfig ...	Lowest priority



Usage: Update default storage options at role level

```
names=> alter role student set
gp_default_storage_options='appendonly=true,compresstype=zlib';
Names=> set role student;
```

Review

- Tablespaces and tablespaces
- Additional table types (external, temp)
- Table storage models
- Compression options
- Test it out in the lab

Pivotal

A NEW PLATFORM FOR A NEW ERA