

# Pivotal

A NEW PLATFORM FOR A NEW ERA

# GPDB Data Loading



Pivotal® Greenplum  
Database


# Agenda

- Introduction
- Data Loading Methods
- Data Loading Performance Tips
- Test it out in the lab


# Data Loading Methods




External tables




gpfdist




gpload



gphdfs



SQL COPY



SQL INSERT

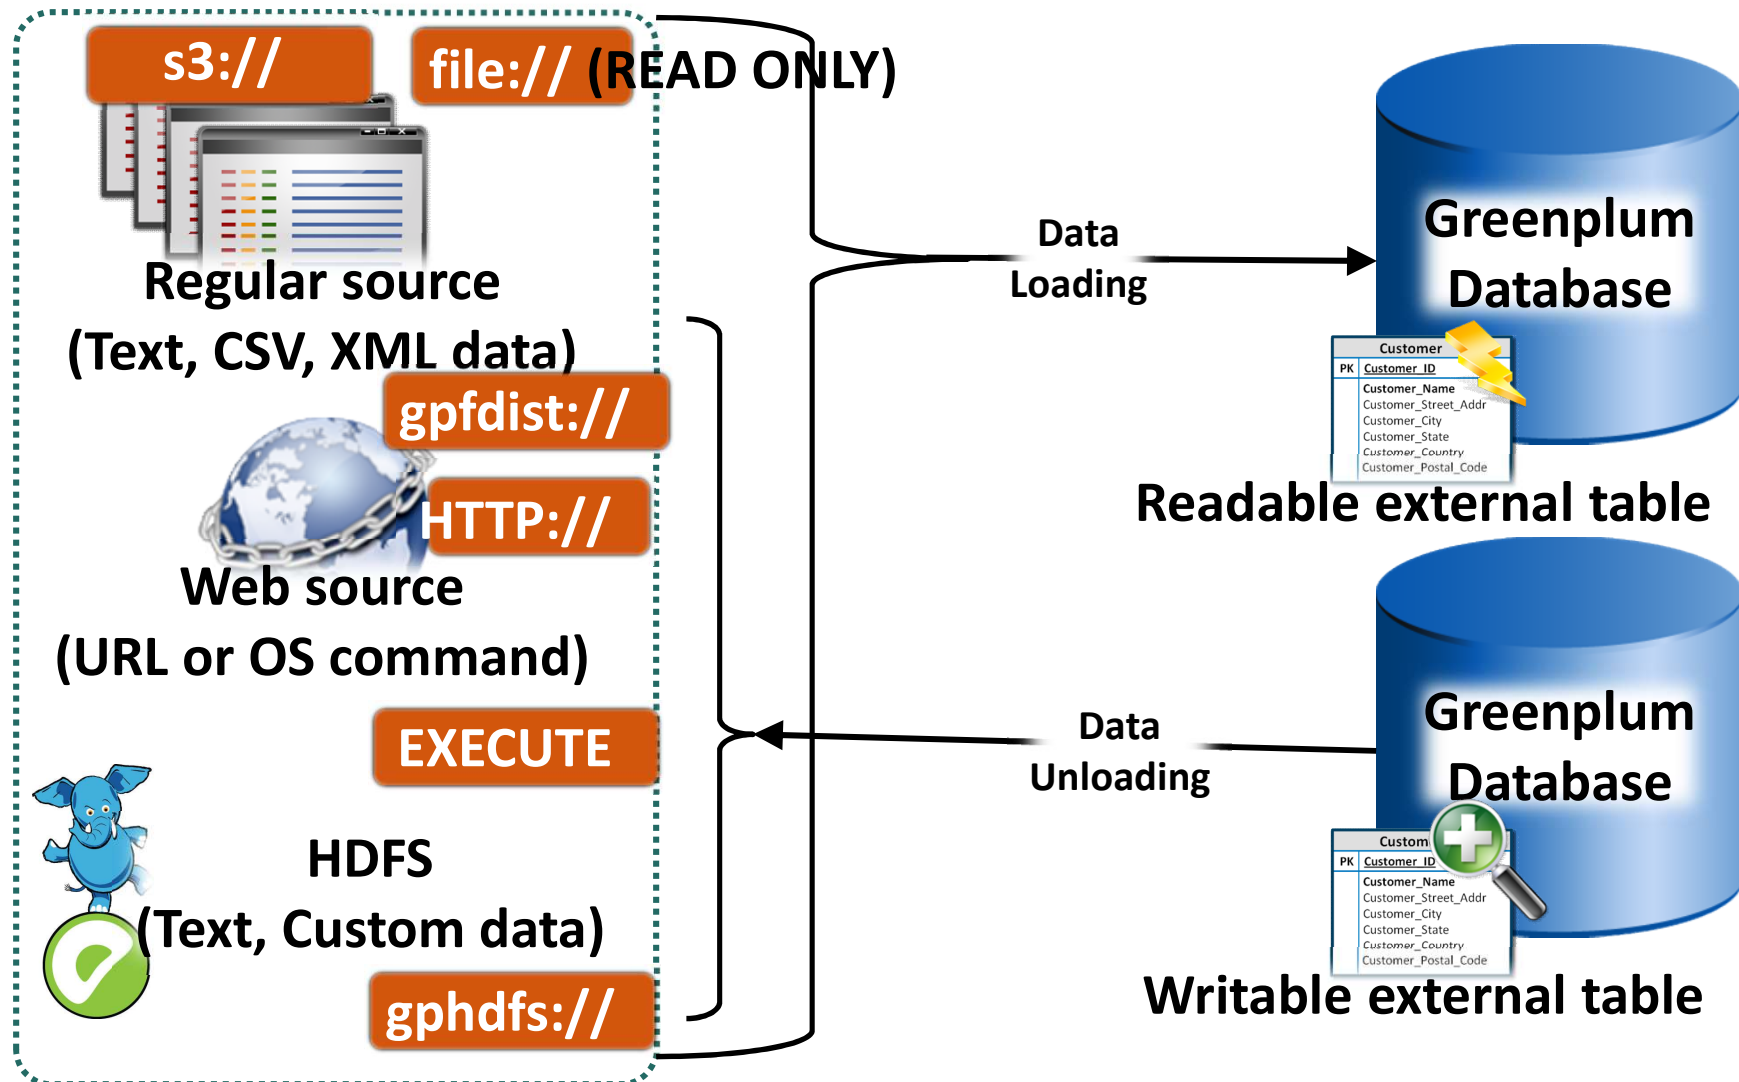
# Loading with External Tables

Read-only external tables:

- Leverage parallel processing power of segments
- Can be accessed with `SELECT` statements
- Access data outside of the Greenplum Database
- Commonly used for ETL and data loading

# External Table Types

## Data Sources and Protocols



# File-Based External Tables

When creating file-based external tables:

- Specify up to as many URIs as you have segments in the `LOCATION` clause
- Each URI points to an external data file or data source
- URIs do not need to exist prior to defining the external table
- The URI must exist when the data is queried
- Great Blog article:  
[File Protocol for External Tables](#)

# File-Based External Table Protocol and Format



## Example: Create an external table with multiple URIs

```
CREATE EXTERNAL TABLE ext_expenses (name text, date date,  
amount float4, category text, description text)  
LOCATION (  
  'file://seghost1/dbfast/external/expenses1.csv',  
  'file://seghost1/dbfast/external/expenses2.csv',  
  'file://seghost2/dbfast/external/expenses3.csv',  
  'file://seghost2/dbfast/external/expenses4.csv',  
  'file://seghost3/dbfast/external/expenses5.csv',  
  'file://seghost3/dbfast/external/expenses6.csv',  
)  
FORMAT 'CSV' ( HEADER );
```

Protocol can be file,  
gpfdist, gpfdists,  
or gphdfs

Format can be  
CSV, TEXT, XML,  
or custom

You can define as  
many URIs as you  
have segments



# Parallel File Distribution Program

The parallel file distribution program, `gpfdist`:

- Is a C program that uses HTTP
- Can be run on an external server
- Distributes data at 200 MB/s per `gpfdist`
- Provides full parallelism for best performance

The data load utility, `gpload`:

- Interfaces with and invokes `gpfdist`
- Creates an external table definition
- Executes `INSERT`, `UPDATE`, or `MERGE` to load data

# gpfdist Based External Tables

When creating file-based external tables:

- Every segment opens a connection to the remote gpfdist agent (try running it with -V)
- Every segment asks for a block of data (ie. just a bunch of data only parsed for EOL)
  - Max block length must be between 32K and 256MB
- After the data arrives at the segment, then the DSITRIBUTED BY clause is invoked which likely will cause data movement to another segment

# Parallel File Distribution Program Example

```
gpfdist -d /var/load_files/expenses1 -p 8081 >> gpfdist.log  
2>&1 &  
gpfdist -d /var/load_files/expenses2 -p 8082 >> gpfdist.log  
2>&1 &
```



## Example: Creating an external table using the gpfdist protocol

```
CREATE EXTERNAL TABLE ext_expenses  
  (name text, date date, amount float4, description text)  
  LOCATION (  
    'gpfdist://etlhost:8081/*',  
    'gpfdist://etlhost:8082/*')  
  FORMAT 'TEXT' (DELIMITER '|')  
  ENCODING 'UTF-8'  
  LOG ERRORS INTO ext_expenses_loadererrors  
  SEGMENT REJECT LIMIT 10000 ROWS ;
```

# Parallel File Distribution Program Example (Cont'd.)



**Example: Load data into a regular table**

```
INSERT INTO expenses (SELECT * FROM ext_expenses);
```

# Accessing Hadoop Data Using gphdfs

```
2 DROP EXTERNAL TABLE IF EXISTS wiki_pages_ext;
3 CREATE EXTERNAL TABLE wiki_pages_ext
4 (
5     LIKE wiki_pages
6 )
7 LOCATION ('gphdfs://hadoop-w-0:8020/user/gpadmin/word_count/out/part-m-*.')
8 FORMAT 'TEXT' (DELIMITER E'\t' NULL E'')
9 ENCODING 'UTF8'
10 LOG ERRORS INTO wiki_err SEGMENT REJECT LIMIT 1 PERCENT;
11 GRANT SELECT ON wiki_pages_ext TO demo;
12 /*
13 demo=# INSERT INTO wiki_pages
14 demo=# SELECT * FROM wiki_pages_ext;
15 NOTICE: Found 66 data formatting errors (66 or more input rows). Rejected 1
16 INSERT 0 15347677
17 Time: 9674.557 ms
18 demo=# select count(*) from wiki_pages;
19      count
20 -----
21 15347677
22 (1 row)
23 */
```

# Accessing Hadoop Data Using gphdfs (Cont'd.)

One time setup for using gphdfs requires setting a couple of configuration parameters:

`gp_hadoop_target_version` -- corresponds with the version of Hadoop you're using

`gp_hadoop_home` -- points to your Hadoop client installation directory

# gpload YAML Control File Example



## Example: YAML Control File for Loading Data

```
VERSION: 1.0.0.1
DATABASE: faa2
USER: gpadmin
HOST: smdw
PORT: 5432
Gpload:
```

Greenplum Database  
connection information

gpload definition  
block

INPUT:

```
- SOURCE:
  LOCAL_HOSTNAME:
    - mdw
  PORT: 8081
  FILE:
    - /rawdata/FAADData/On_Time_On_Time_Performance_*.csv
- FORMAT: csv
- DELIMITER: ','
- ESCAPE: '"'
- QUOTE: '"'
- HEADER: true
- ERROR_LIMIT: 100
- ERROR_TABLE: faadata.faadataerr
```

OUTPUT:

```
- TABLE: faadata.factontimeperformance
- MODE: insert
```

PRELOAD:

```
- TRUNCATE: true
- REUSE_TABLES: false
```

SQL:

```
- BEFORE: "INSERT INTO audit VALUES('start', current_timestamp)"
- AFTER: "INSERT INTO audit VALUES('end', current_timestamp)"
```

gpfdist  
configuration  
information

Pre-existing table  
and data entry  
mode

Pre and post  
SQL statements

Indentation  
spaces, **not**  
**tabs**,  
determine  
hierarchy

Preloading  
data section

# gpload Syntax

The gpload syntax is as follows:

```
gpload -f control_file [-l log_file] [-h  
hostname] [-p port] [-U username] [-d  
database] [-W] [-v | -V] [-q] [-D] gpload  
-? | --version
```

The following is an example of how it is used:



## Example: Successful data load with gpload

```
$ gpload -f load_faadata.yaml
2012-01-17 09:05:29|INFO|gpload session started 2012-01-17 09:05:29
2012-01-17 09:05:29|INFO|started gpfdist -p 8081 -P 8082 -f
"/rawdata/FAADOn_Time_Performance_*.csv" -t 30
2012-01-17 09:11:23|INFO|running time: 353.36 seconds
2012-01-17 09:11:23|INFO|rows Inserted          = 20860045
2012-01-17 09:11:23|INFO|rows Updated          = 0
2012-01-17 09:11:23|INFO|data formatting errors = 0
2012-01-17 09:11:23|INFO|gpload succeeded
$
```



# Loading Data with gpload



## Example: Load data with gpload

```
$ gpload -f load_faadata.yaml
```

gpadmin@mdw:~/course\_samples

```
[gpadmin@mdw course_samples]$ gpload -f load_faadata.yaml
2012-01-19 06:10:57 |INFO| gpload session started 2012-01-19 06:10:57
2012-01-19 06:10:57 |INFO| started gpfdist -p 8081 -P 8085 -f "/rawdata/FAADData/On
_Time_On_Time_Performance_*.csv" -t 30
```

gpload starts gpfdist

### Query Text

```
INSERT INTO "faadata"."factontimeperformance"
("year","quarterid","monthid","dayofmonth","dayid","flightdate","tailnum","flightnum","originairportid","origincityname","originstateabbreviation","originstatefipscode","originstatename","originwacid","destairportid","destcityname","deststateabbreviation","deststatefipscode","deststatename","destwacid","crsdeptime","deptime","depdelay","depdelay2","depdel15","deptimeblkid","taxiout","wheelsoff","wheelson","taxiin","crsarrrtime","arrtime","arrdelay","arrdelay2","arrdel15","arrdelaygroupid","arrtimeblkid","cancelled","cancellationid","diverted","crselapsedtime","actualelapsedtime","airtime","numflights","distance","distancegroupid","carrierdelay","weatherdelay","nasdelay","securitydelay","lateaircraftdelay","firstdeptime","totaladdgtime","longestaddgtime","divairportlandings","divertedreacheddest","divactualelapsedtime","divarrdelay","divdistance","div1airport","div1wheelson","div1totalgtime","div1longestgtime","div1wheelsoff","div1tailnum","div2airport","div2wheelson","div2totalgtime","div2longestgtime","div2wheelsoff","div2tailnum","div3airport","div3wheelson","div3totalgtime","div3longestgtime","div3wheelsoff","div3tailnum","div4airport","div4wheelson","div4totalgtime","div4longestgtime","div4wheelsoff","div4tailnum","div5airport","div5wheelson","div5totalgtime","div5longestgtime","div5wheelsoff","div5tailnum") SELECT
"year","quarterid","monthid","dayofmonth","dayid","flightdate","uniqucarrierid","airlineid","carrierid","tailnum","flightnum","originairportid","origincityname","originstateabbreviation","originstatefipscode","originstatename","originwacid","destairportid","destcityname","deststateabbreviation","deststatefipscode","deststatename","destwacid","crsdeptime","deptime","depdelay","depdelay2","depdel15","deptimeblkid","taxiout","wheelsoff","wheelson","taxiin","crsarrrtime","arrtime","arrdelay","arrdelay2","arrdel15","arrdelaygroupid","arrtimeblkid","cancelled","cancellationid","diverted","crselapsedtime","actualelapsedtime","airtime","numflights","distance","distancegroupid","carrierdelay","weatherdelay","nasdelay","securitydelay","lateaircraftdelay","firstdeptime","totaladdgtime","longestaddgtime","divairportlandings","divertedreacheddest","divactualelapsedtime","divarrdelay","divdistance","div1airport","div1wheelson","div1totalgtime","div1longestgtime","div1wheelsoff","div1tailnum","div2airport","div2wheelson","div2totalgtime","div2longestgtime","div2wheelsoff","div2tailnum","div3airport","div3wheelson","div3totalgtime","div3longestgtime","div3wheelsoff","div3tailnum","div4airport","div4wheelson","div4totalgtime","div4longestgtime","div4wheelsoff","div4tailnum","div5airport","div5wheelson","div5totalgtime","div5longestgtime","div5wheelsoff","div5tailnum" FROM ext_gpload20120119_060258_15300
0 appname
```

Data is loaded into target table

# External Web Table Protocols and Format



## Example: External WEB table with data loads from URLs

```
CREATE EXTERNAL WEB TABLE ext_expenses (name text, date
date, amount float4, description text)
  LOCATION (
    'http://intranet.company.com/expenses/sales/expenses.csv',
    'http://intranet.company.com/expenses/finance/expenses.csv',
    'http://intranet.company.com/expenses/ops/expenses.csv')
  FORMAT 'CSV' ( HEADER );
```



## Example: External WEB table with data loads from a script

```
CREATE EXTERNAL WEB TABLE log_output (linenum int, message
text)
  EXECUTE '/var/load_scripts/get_log_data.sh'
  ON HOST FORMAT 'TEXT' (DELIMITER '|');
```

# External Web Table Protocols and Format (Cont'd.)



## Example: External WEB table with data loads from a command

```
CREATE EXTERNAL WEB TABLE du_space (storage text)
EXECUTE 'df -k' ON HOST FORMAT 'TEXT';
```

# Environment Variables for Command-Based Web Tables

- Are not sourced at the segment host
- Can be set in the EXECUTE clause as follows:



## Example: External WEB table with environment variables

```
CREATE EXTERNAL WEB TABLE TEST(segname text, abbrev text)
EXECUTE 'export HNAME="$HOSTNAME"-SR; echo
"$HOSTNAME,$HNAME"' FORMAT 'TEXT' (DELIMITER ',');
SELECT * FROM test;
segname | abbrev
-----+-----
sdw2    | sdw2-SR
sdw1    | sdw1-SR
(2 rows)
```



**Note:** You can disable the use of the EXECUTE command in web table definitions by setting `gp_external_enable_exec` to off.

# External Table Environment Variables

Variable	Description
<code>\$GP_CID</code>	Command count of the session executing the external table statement.
<code>\$GP_DATABASE</code>	The database that the external table definition resides in.
<code>\$GP_DATE</code>	The date the external table command was executed.
<code>\$GP_MASTER_HOST</code>	The host name of the Greenplum master host from which the external table statement was dispatched.
<code>\$GP_MASTER_PORT</code>	The port number of the Greenplum master instance from which the external table statement was dispatched.
<code>\$GP_SEG_DATADIR</code>	The location of the data directory of the segment instance executing the external table command.
<code>\$GP_SEG_PG_CONF</code>	The location of the <code>postgresql.conf</code> file of the segment instance executing the external command.

# External Table Environment Variables (cont'd.)

Variable	Description
<code>\$GP_SEG_PORT</code>	The port number of the segment instance executing the external table command.
<code>\$GP_SEGMENT_COUNT</code>	The total number of primary segment instances in the Greenplum Database system.
<code>\$GP_SEGMENT_ID</code>	The ID number of the segment instance executing the external table command.
<code>\$GP_SESSION_ID</code>	The database session identifier number associated with the external table statement.
<code>\$GP_SN</code>	Serial number of the external table scan node in the query plan of the external table statement.
<code>\$GP_TIME</code>	The time the external table command was executed.
<code>\$GP_USER</code>	The database user executing the external table statement.
<code>\$GP_XID</code>	The transaction ID of the external table statement.

# External Table Error Handling

When handling errors using external tables:

- Incorrectly formatted rows are rejected:
  - Rows with missing or extra attributes
  - Rows with columns of the wrong data type
  - Rows with invalid client encoding sequence
- CONSTRAINT errors (NOT NULL, CHECK, UNIQUE ) are handled “all or nothing” -- not single row isolation
- Format of error handling clause:  

```
[LOG ERRORS INTO error_table] SEGMENT  
REJECT LIMIT count [ROWS | PERCENT]
```

# External Tables and Planner Statistics

Query planning of complex queries on external tables is not optimal because:

- Data resides outside the database
- No database statistics exist for external table data
- Data from external tables are not meant for frequent or ad-hoc access



# COPY SQL Command

The COPY SQL command:

- Is a PostgreSQL command
- Loads all rows in one command and is not parallel
- Loads data from a file or from standard input
- Supports error handling similar to external tables

The following is an example of the command:



**Example: Copy data from `/data/myfile.csv` into the table specified**

```
COPY mytable FROM '/data/myfile.csv' WITH CSV HEADER;
```

# Data Loading Performance Tips

- Drop indexes and recreate them after loading data
- Use gpfdist to load or unload data in Greenplum Database
- Spread the data evenly across as many ETL nodes as possible
- Split very large data files into equal parts and spread the data across as many file systems as possible
- Run two gpfdist instances per file system
- Run gpfdist on as many network interfaces as possible

## Data Loading Performance Tips (Cont'd.)

- Use `gp_external_max_segs` to control the number of segments each `gpfdist` serves
- Always keep `gp_external_max_segs` and the number of `gpfdist` processes an even factor
- Always drop indexes before loading into existing tables and re-create the index after loading
- Always run `ANALYZE` on the table after loading it
- Disable automatic statistics collection during loading by setting `gp_autostats_mode` to `NONE`
- Run `VACUUM` after load errors to recover space

# Review

- Data Loading Methods
- Data Loading Performance Tips
- Test it out in the lab

# Pivotal

A NEW PLATFORM FOR A NEW ERA