

Pivotal

A NEW PLATFORM FOR A NEW ERA

Physical Design Considerations in GPDB



Pivotal® Greenplum
Database

Pivotal™

Agenda

- Introduction
- Identify optimal distribution key
- Check for data skew
- Table partitioning: benefits, when to partition
- Select appropriate data types
- Define constraints
- Put this into practice in the lab

Key Design Considerations

The following are key design considerations to account for:



Data distribution and distribution key selection



Checking for data skew



To partition or not



Choosing appropriate data types



Defining constraints

Data Distribution

In defining distribution methods, there are two options:

- Column distribution with the DISTRIBUTED BY clause
- Random distribution with the DISTRIBUTED RANDOMLY clause

```
CREATE TABLE  tablename  (
column_name1      data_type NOT NULL,
column_name2      data_type NOT NULL,
column_name3      data_type NOT NULL DEFAULT default_value,
. . . )
[DISTRIBUTED BY (column_name)]          hash algorithm
[DISTRIBUTED RANDOMLY]                 random algorithm
```



Note: Every table in the Greenplum database has a distribution method, whether you select it or not.

Distribution Key Considerations

Consider the following:

- A distribution key can be modified, even after the table has been created.
- If a table has a unique constraint, it must be declared as the distribution key.
- User defined data types and geometric data types are not eligible as distribution keys.



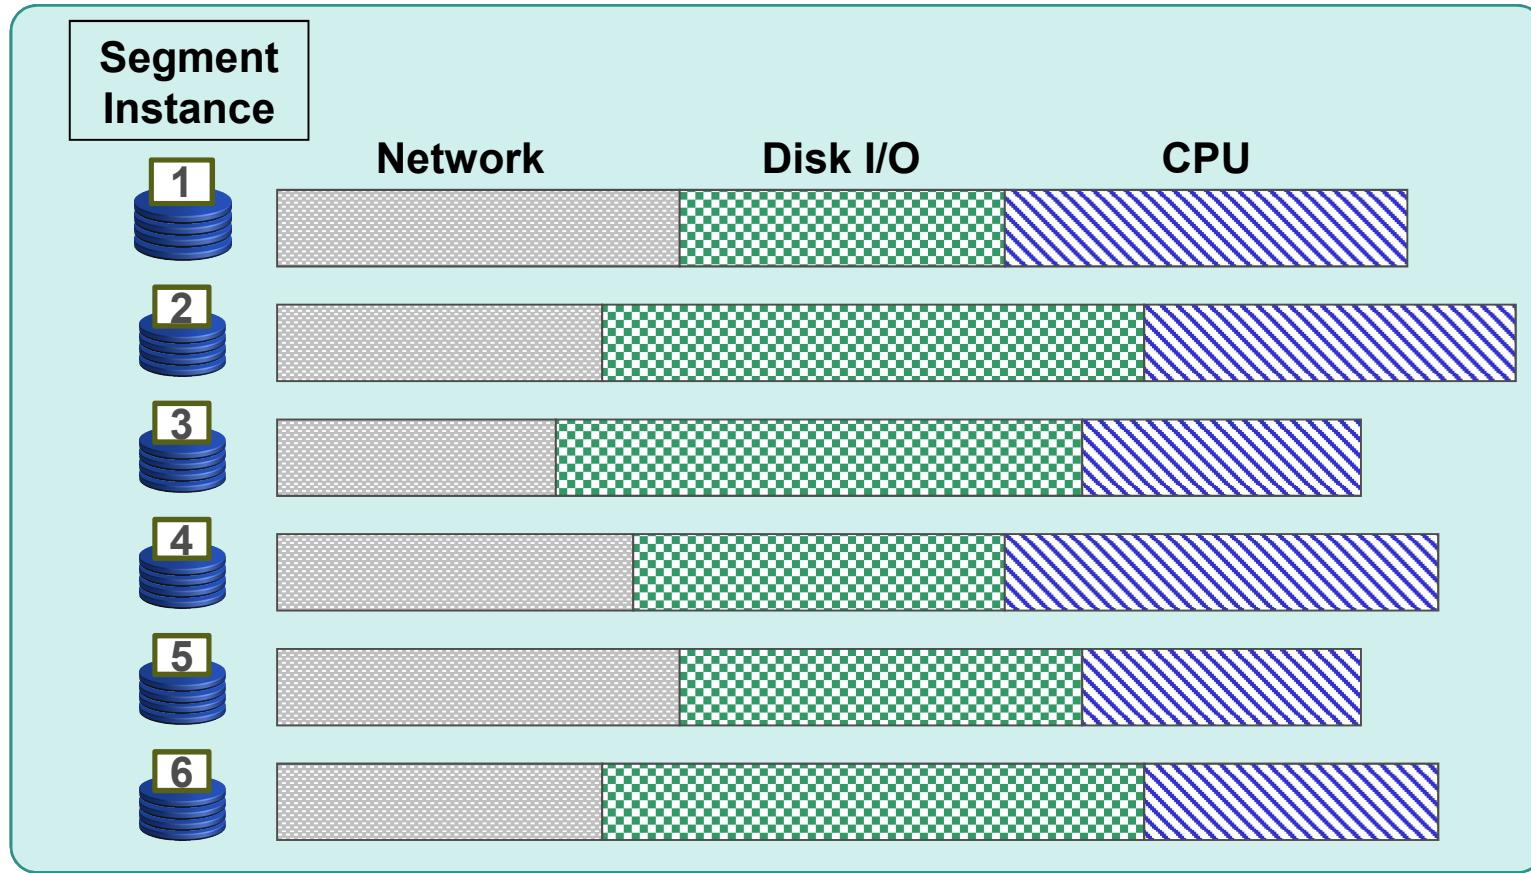
Note: You should choose a distribution key with high cardinality

Default Distribution Use

Consider the following when creating tables:

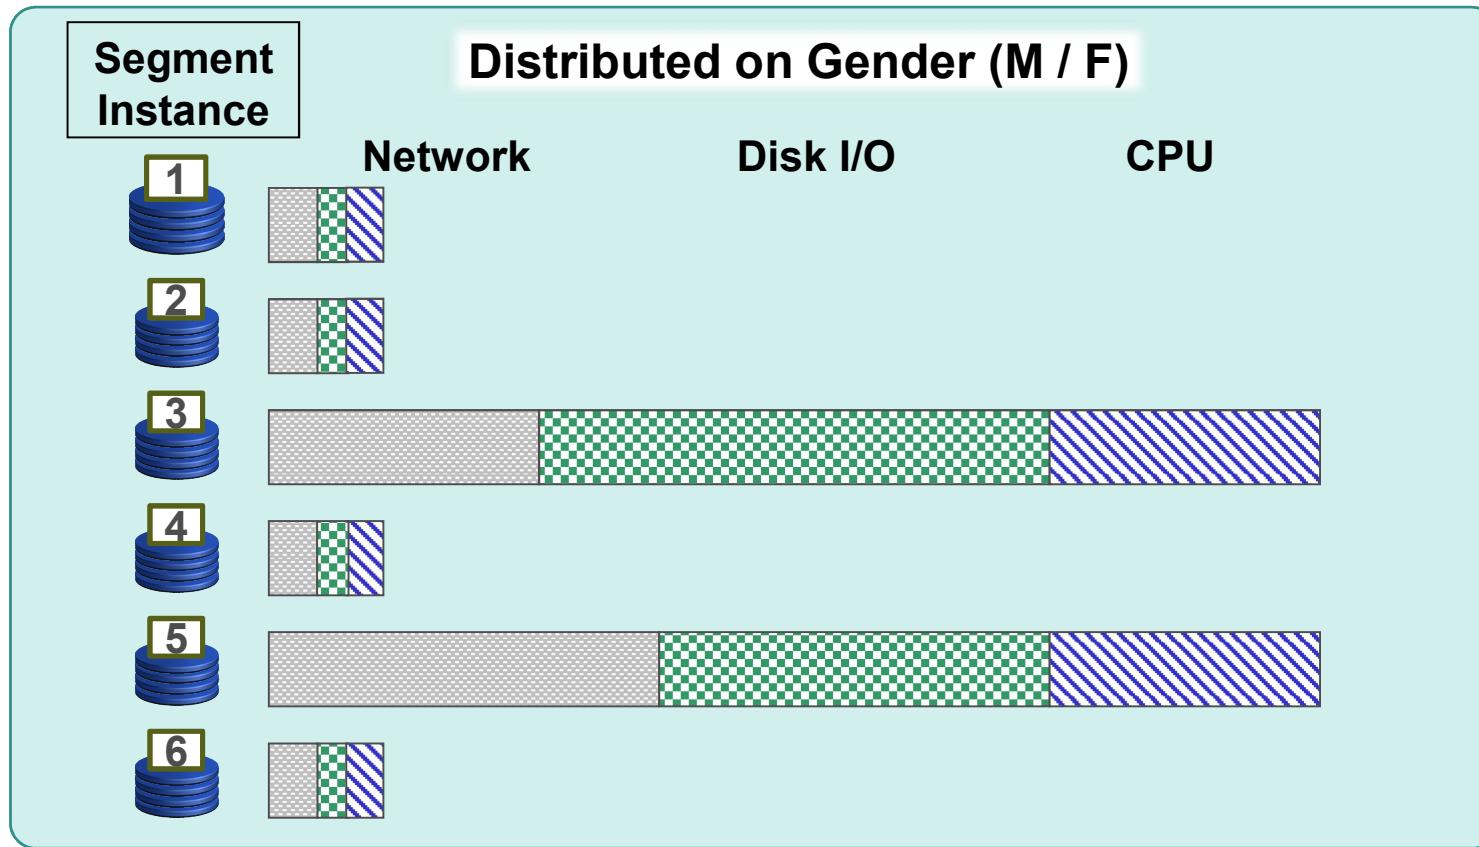
- If a table has a primary key and a distribution key is not specified, by default the primary key will be used as the distribution key.
- If the table does not have a primary key and a distribution key is not specified, then the first eligible column of the table will be used as the distribution key.
- User defined data types and geometric data types are not eligible as distribution keys.
- If a table does not have an eligible column then a random distribution is used.

Distributions and Performance



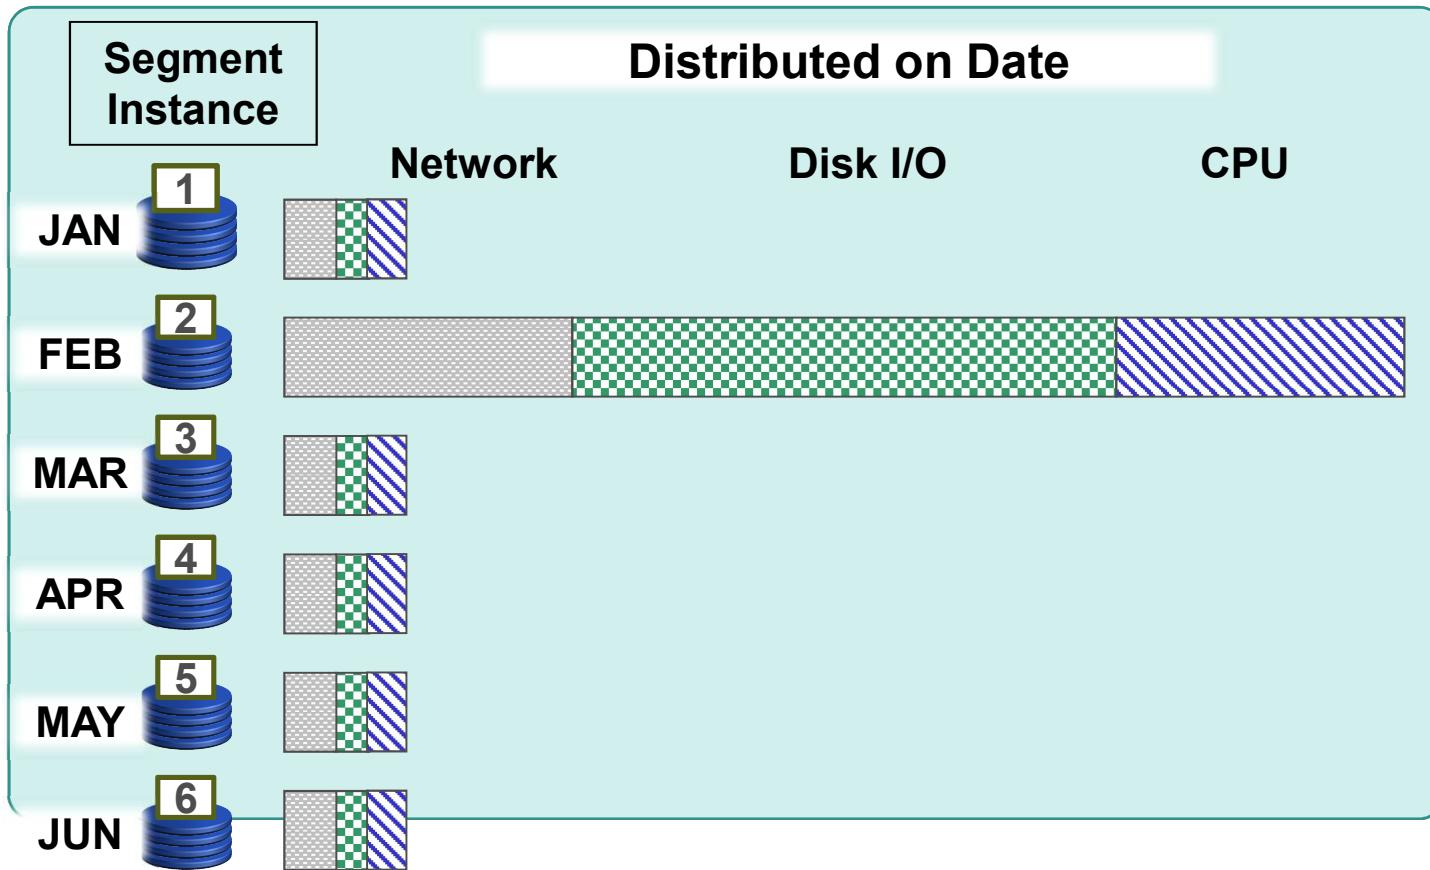
Note: To optimize performance, use a hash distribution method (DISTRIBUTED BY) that distributes data evenly across all segment instances.

Hash Distributions and Data Skew



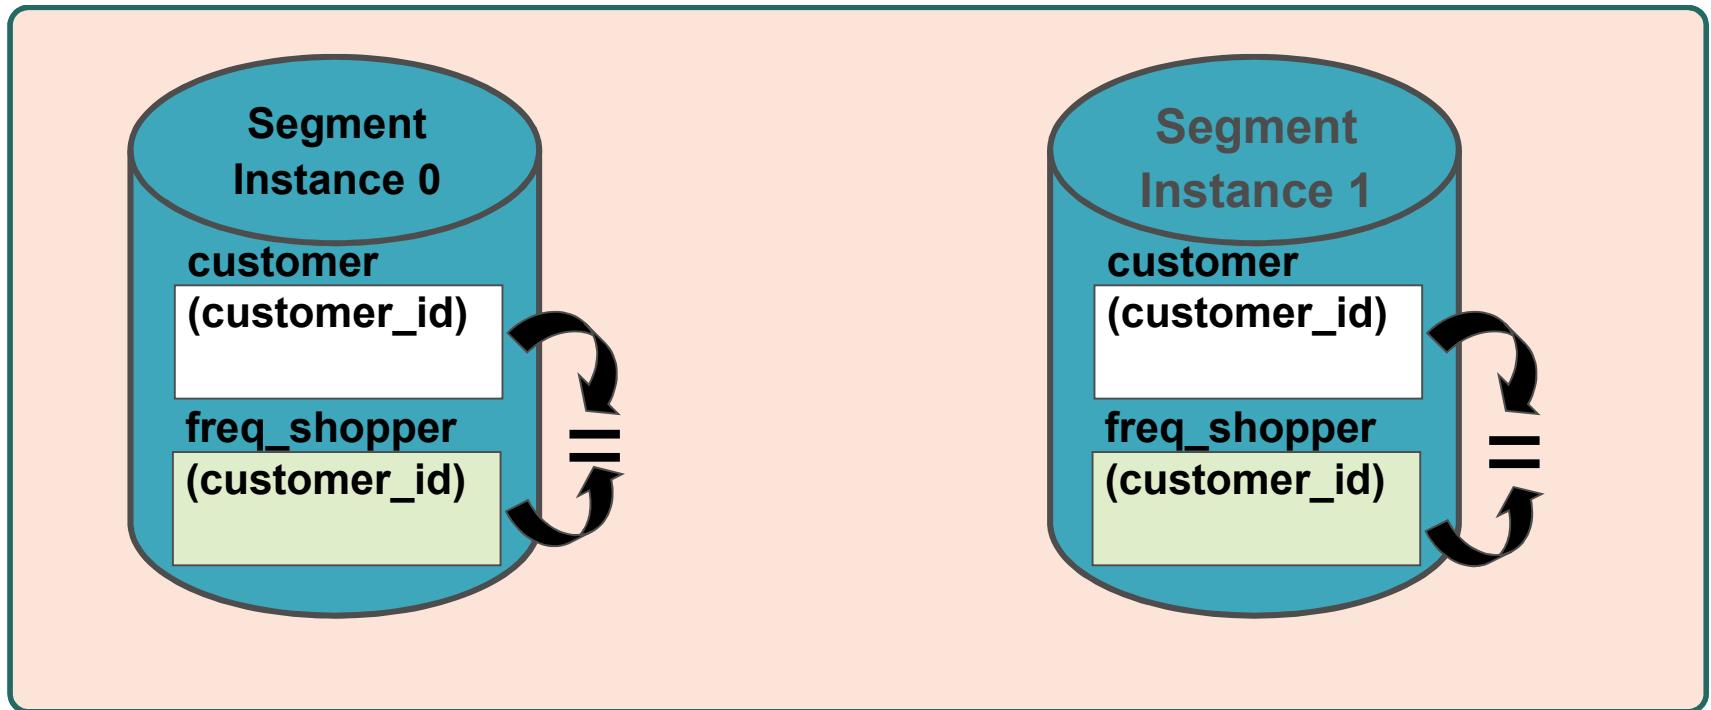
Note: Select a distribution key with unique values and high cardinality.

Hash Distributions and Processing Skew



Note: Select a distribution key that will not result in processing skew.

Using the Same Distribution Key for Commonly Joined Tables



Note: Optimize for local joins. Distribute on the same key used in the JOIN.

Caveat to Previous Slide: ***Use the Same Data Type*** for These Distribution Keys

If you use *different data types for identical values*:

- Different types, containing identical values, hash to different values
 - Resulting in like rows being stored on different segment instances
 - Requiring a redistribution before the tables can be joined

The following shows two distribution keys with different data types:

```
customer (customer_id)    745::int
freq_shopper (customer_id) 745::varchar(10)
```

Distributed With DISTRIBUTED RANDOMLY

The DISTRIBUTED RANDOMLY clause:

- Uses a random algorithm
- Requires a redistribute or broadcast motion for any query that joins to a table having this distribution policy
- Is acceptable for small tables such as dimension tables
- Is not acceptable for large tables such as fact tables

Distribution Best Practices

- Explicitly define a column or random distribution for all tables. Do not use the default.
- Use a single column that will distribute data across all segments evenly.
- Do not distribute on columns that will be used in the WHERE clause of a query.
- Do not distribute on dates or timestamps.
- Never distribute and partition tables on the same column.
- Achieve local joins to significantly improve performance by distributing on the same column for large tables commonly joined together.
- Validate that data is evenly distributed after the initial load and after incremental loads.

Check for Data Skew

Check for data skew with the following:

- gp_toolkit administrative schema offers two views:
 - gp_toolkit(gp_skew_coefficients)
 - gp_toolkit(gp_skew_idle_fractions)
- To view the number of rows on each segment, run the following query:

```
SELECT gp_segment_id, count(*)  
FROM table_name GROUP BY gp_segment_id;
```

- Check for processing skew with the following query:

```
SELECT gp_segment_id, count(*) FROM table_name  
WHERE column='value' GROUP BY gp_segment_id;
```

Processing and Data Skew Example



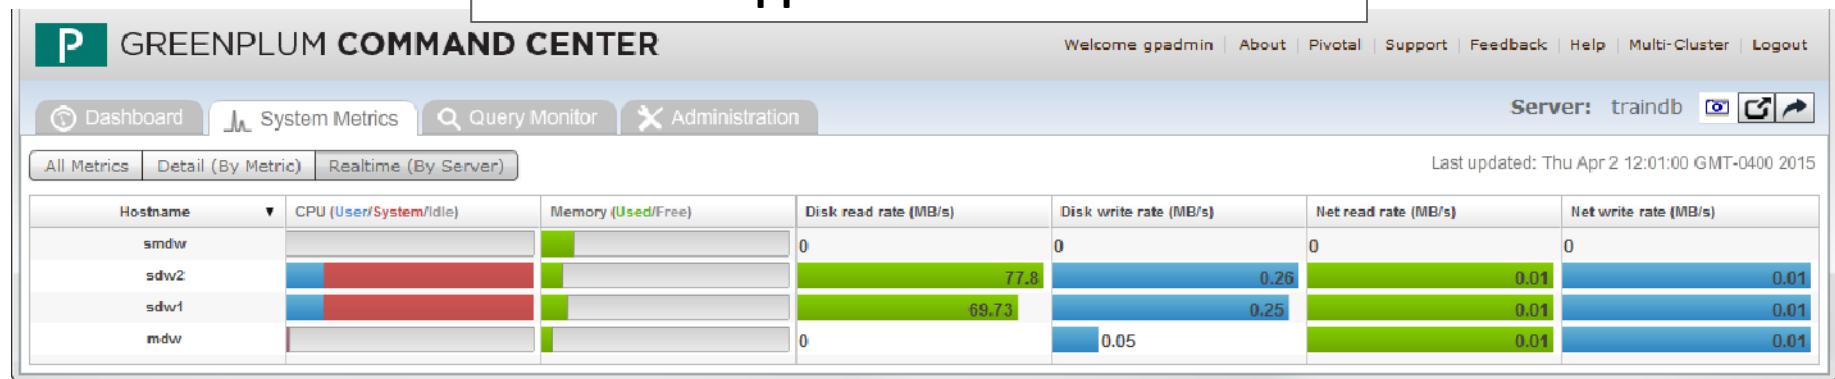
Example: Determining the processing and data skew of existing tables

```
SELECT sifrelename, siffraction, skccoeff
from gp_skew_idle_fractions, gp_skew_coefficients
WHERE sifrelename=skcrelname AND siffraction > 0.1;
```

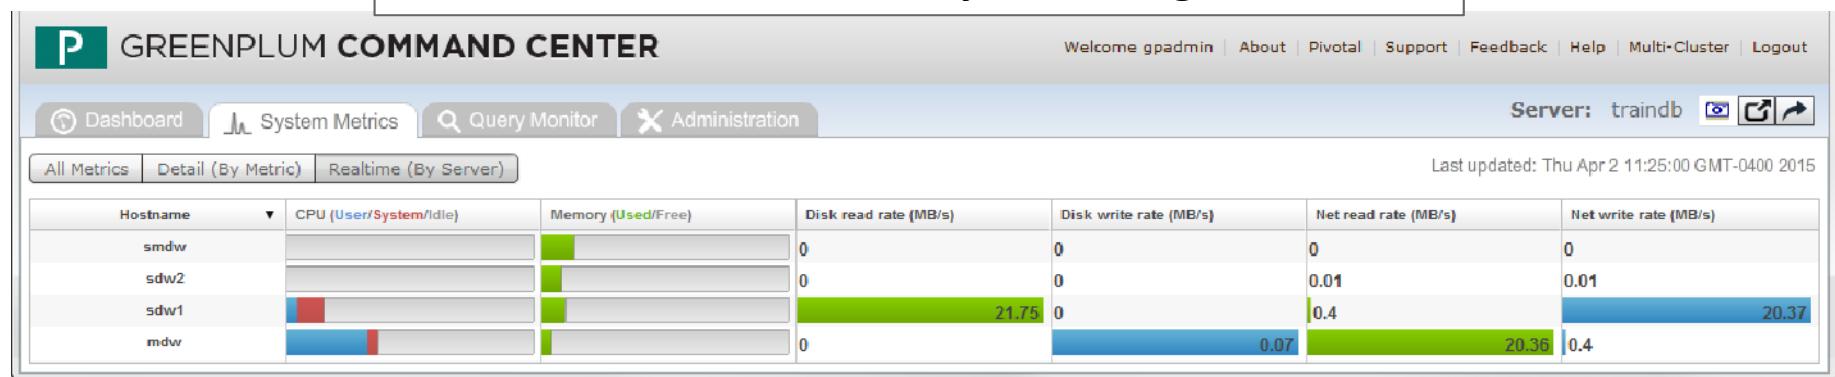
sifrelename	siffraction	skccoeff
car_data	0.5000000000000000	141.42135623730949000
correlation	0.4000000000000000	94.280904158206336667000
golfevaluate	0.2222222222222222	40.406101782088430000000
golfnew	0.2222222222222222	40.406101782088430000000
golftest	0.2222222222222222	40.406101782088430000000
old_regr_example	0.5000000000000000	141.42135623730950000
old_regr_example_test	0.5000000000000000	141.42135623730951000
sales_example	0.5000000000000000	141.42135623730951000
water_treatment_predict	0.12037037037037037	19.352396116684458526000
winequality_white	0.10293040293040293040	16.226786893705185790000

Real-Time Data and Processing Skew in Command Center

Data access appears to be even distributed



Data access is skewed heavily to one segment server



Redisistribute Using ALTER TABLE

Use the ALTER TABLE command to:

- Redisistribute on the *current* policy

```
ALTER TABLE sales SET WITH (REORGANIZE=TRUE);
```

- Redisistribute with a *new* distribution key

```
ALTER TABLE sales SET DISTRIBUTED BY (column);
```

Why Partition a Table?

The following are reasons to partition a table:

- Provide more efficiency in querying a large table, by using partition elimination to avoid full table scans
- Without the overhead and maintenance costs of an index on a date column, for example
- To handle increasing volumes of data that are not needed by the average query
- Allow instantaneous dropping of older data and simple addition of newer data
- Supports a rolling n period methodology for transactional data

Candidates for Partitioning

The following date-related tables make excellent candidates for table partitioning:

- Fact tables with dates
- Fact tables based on region (e.g. "EMEA", "APAC", ...)
- Transaction tables with dates
- Activity tables with dates
- Statement tables with numeric statement periods (YYYYMM)
- Financial summary tables with end of month dates

Table Partitioning Best Practices

Use table partitioning:

- On large distributed tables to improve query performance.
- If the table can be divided into rather equal parts based on a defining criteria, such as by date.
- If the defining partitioning criteria is the same access pattern used in query predicates, such as
`WHERE txn_date = '2015-05-14'`
- If possible, avoid using default partitions in your design (they are *always* scanned).

Creating Partitioned Tables

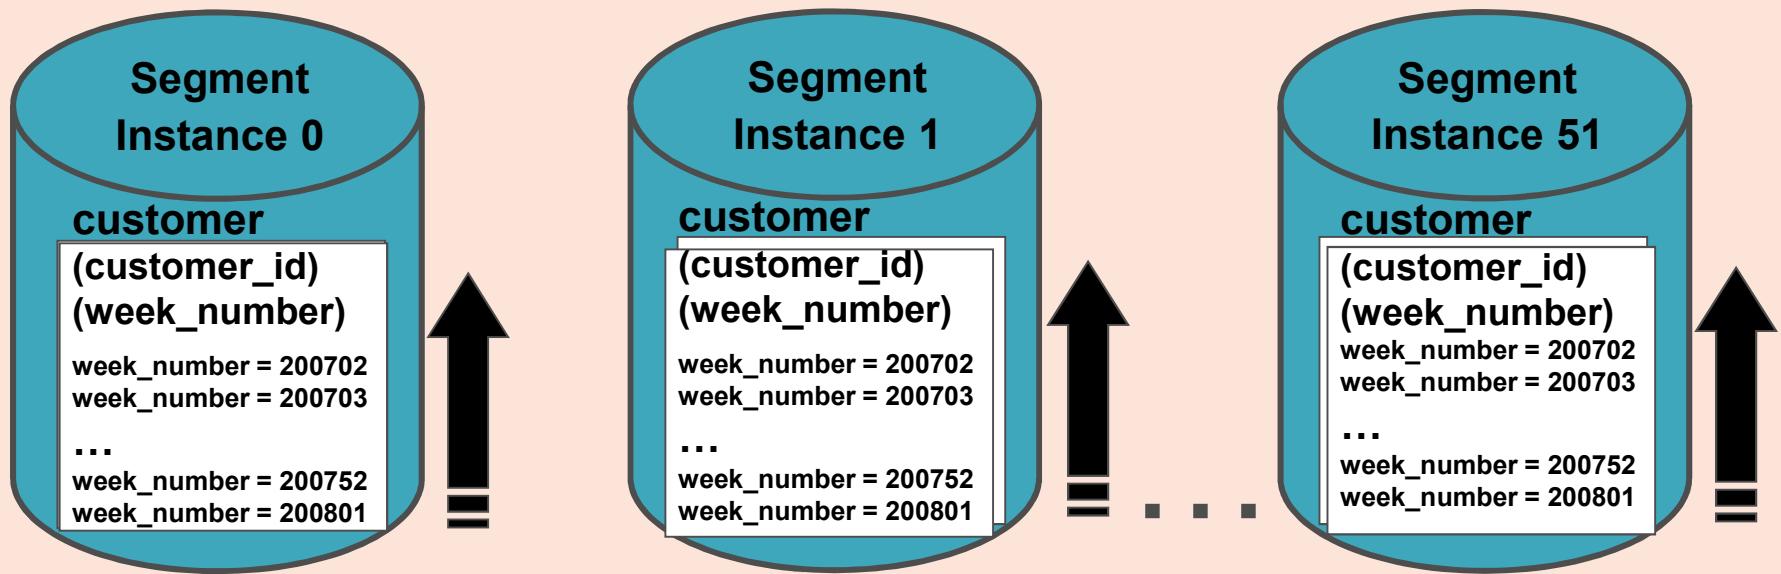
The following shows how to create the customer partition table using an interval:

Example: Create a partition table using an interval

```
CREATE TABLE customer (
    customer_id      INT,
    week_number      INT,
    . . .
) DISTRIBUTED BY (customer_id)
PARTITION BY RANGE ( week_number )
(START (200701) END (200752) INCLUSIVE EVERY (1));
```

Maintaining Rolling Windows

Use ALTER TABLE with ADD PARTITION clause to add a new child table



Use ALTER TABLE with DROP PARTITION clause to delete an old child table

Data Type Best Practices

The following should be considered when choosing data types:

- Use data types to constrain your column data:
 - Use character types to store strings
 - Use date or timestamp types to store dates
 - Use numeric types to store numbers
- Choose the type that uses the least space:
 - Do not use BIGINT when INTEGER is sufficient
 - Use identical data types for columns used in join operations
 - Converting data types prior to joining leads to unnecessary data movement

Table and Column Constraints

The following code creates:

- A table with a CHECK constraint:

```
CREATE TABLE products
( product_no integer, name text,
  price numeric CHECK (price >= 0) );
```

- A table with a NOT NULL constraint:

```
CREATE TABLE products
( product_no integer NOT NULL, name text NOT NULL,
  price numeric );
```

Table and Column Constraints (Cont)

- A table with a **UNIQUE** constraint:

```
CREATE TABLE products
( product_no integer UNIQUE, name text, price numeric)
DISTRIBUTED BY (product_no);
```

- A table with a **PRIMARY KEY** constraint

```
CREATE TABLE products
( product_no integer PRIMARY KEY, name text, price
numeric)
DISTRIBUTED BY (product_no);
```

Primary Key and the Distribution Key

Consider the following:

- A primary key is a logical model concept which allows each row to be uniquely identified.
- A DISTRIBUTED BY key is a physical Greenplum Database concept which determines where, within the cluster, each row is physically stored.
- The distributed by key can differ from the primary key, as long as it is a *left-subset* of the primary key.
- This applies to UNIQUE constraints as well.

Review

- Identify optimal distribution key
- Check for data skew
- Table partitioning: benefits, when to partition
- Select appropriate data types
- Define constraints
- Put this into practice in the lab

Pivotal

A NEW PLATFORM FOR A NEW ERA