

# Pivotal

A NEW PLATFORM FOR A NEW ERA

# SQL Joins: Types and Implementation








Pivotal® **Greenplum**  
**Database**

Pivotal

# Agenda


- What Join Types Are Available to You?
- Row Elimination
- Minimize Data Movement
- Join Implementations Seen in Query Plans
- Try the Lab

# JOIN Types



# Inner Join

- Can be simple to write
- Require a join condition be specified
- Are also known as a SIMPLE JOIN or EQUI-JOIN




```
-- Inner join
SELECT c.id, c.name, o.value FROM clients c, orders o
WHERE c.id = o.id
ORDER BY c.id ASC;
```

# ANSI Syntax for Joins

```
joindb=# select * from clients join orders on  
clients.id=orders.id;
```

id	name	id	value
1	bob	1	val1
2	alice	2	val2

(2 rows)




# Left Outer Join

```
joindb=# select * from clients left join orders on  
clients.id=orders.id;
```

id	name	id	value
0	joe		
1	bob	1	val1
2	alice	2	val2

(3 rows)




# Right Outer Join

```
joindb=# select * from clients right join orders on  
clients.id=orders.id;
```

id	name	id	value
1	bob	1	val1
2	alice	2	val2
		3	val3

(3 rows)






# Full Outer Join

```
joindb=# select * from clients full join orders on  
clients.id=orders.id;
```

id	name	id	value
0	joe		
1	bob	1	val1
2	alice	2	val2
		3	val3

(4 rows)



# Natural Full Outer Join


```
joindb=# select * from clients natural full join orders;
```

id	name	value
0	joe	
1	bob	val1
2	alice	val2
3		val3

(4 rows)

# Cross Join (Cartesian Product)

- Often, this is used unintentionally
- Combines every row in the left table with every row in the right table
- Specified with the CROSS JOIN syntax or by comma separating the table names, with *no predicates*




*	id	name	id	value
1	0	Joe	1	val1
2	0	Joe	3	val3
3	0	Joe	2	val2
4	1	Bob	2	val2
5	1	Bob	3	val3
6	1	Bob	1	val1
7	2	Alice	1	val1
8	2	Alice	3	val3
9	2	Alice	2	val2

# Row Elimination During Joins

- Greenplum often has to use disk space to temporarily store data.
- Query optimizer minimizes the amount of memory and disk space required by:
  - Projecting (copying) only those columns that the query requires
  - Doing single-table set selections first (qualifying rows)
  - Eliminating rows early

# Analyzing Row Elimination

- Copies selected rows into temporary tables or spill files
- Projects needed columns into spill file
- Restricts data starting from the `WHERE` clause, then to the `FROM` statement, and finally to the `SELECT` statement



- Temp space consumed is a function of the number of rows times the number of columns

# Joins: Parallel Implementation

- The core join algorithms will be the same as in non-distributed systems.
- Additional details: how do we partition data and still guarantee correctness in a distributed system?
- For the most part, these details are the same, regardless of the join algorithm (e.g. Merge, Hash, Nested Loop).

# Row Redistribution / Motion During Joins

- Generally speaking, if you can do a **co-located join**, you do.
- Between broadcast and redistributed joins, the optimizer looks at the cost of motion that gets introduced and chooses **the plan that it believes to be cheaper**.
- The optimizer depends on up-to-date statistics on the tables.

# Co-located Join

- If the join key is the distribution key for both tables:
- We can guarantee<sup>1</sup> the join can be handled locally on each segment without redistribution.
- This is the most efficient option.
- This is a key consideration when designing your schema.
- *All equal keys are hashed to the same node, so join results will be correct.*

*[1] This assumes that the datatypes of the join keys is the same so that the hash algorithm for the join keys will hash the keys to the same segment.*



# Broadcast Join

- Neither table is distributed by the join keys.
- One table is considerably smaller than the other table.
- We may broadcast the smaller table.
- All segments see **a complete copy** of the smaller table.
- They perform a co-located join.
- *Each row of the larger table can see all target rows in the smaller table, so join results will be correct.*

# Redistributed Join

- One, or both, of the tables are not distributed on the join keys.
- We can redistribute the table(s).
- They perform a co-located join.
- *All equal keys are hashed to the same node, so join results will be correct.*

# Join Implementations

- The means that GPDB uses to join two tables
- These are what you will see in query plans:
  - Sort Merge
  - Hash
  - Nested Loop

# Sort Merge Join

- Common when the join condition is based an inequality operator, like  $<$ ,  $<=$ ,  $>$ ,  $>=$  (**but not**  $<>$ )
- Steps:
  - Sort the tables by the join attribute.
  - Scan the two tables in parallel.
  - Combine matching rows to form join rows.

# Illustration of the Sort Merge Join Process

```
3  Relation output = new Relation();
4  while (!left.IsPastEnd() && !right.IsPastEnd())
5  {
6      if (left.Key == right.Key) {
7          output.Add(left.Key);
8          left.Advance();
9          right.Advance();
10     } else if (left.Key < right.Key) {
11         left.Advance();
12     } else {
13         right.Advance();
14     }
15     return output;
16 }
17
```

# Hash Join

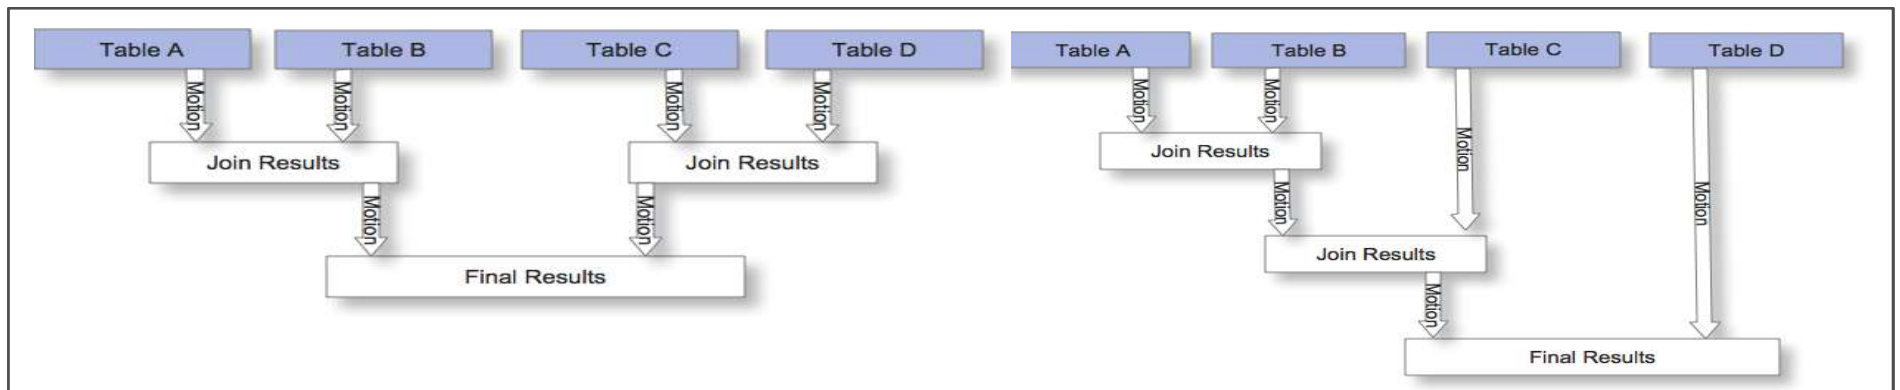
- Build phase:
  - Scan smaller table, creating in-memory hash table.
  - Keys are join columns
  - Values are the corresponding rows
- Probe phase:
  - Scan the larger table and find the relevant rows from the smaller relation by looking in the hash table.

# Nested Loop Join

- Can be one of the most efficient types of joins.
- For each row of the “left table” (loop):
  - Scan the right table to find a match (nested loop)
- Indexes on the join columns make this more efficient
- This type of join is typically inefficient, though

# N-Tables

- All  $n$ -table joins are reduced to a series of two-table joins
- The query engine can only work on two tables at a time
- The final result must be built up by a tree of join steps, each with two inputs.





# Wrap Up

- What Join Types Are Available to You?
- Row Elimination
- Minimize Data Movement
- Join Implementations Seen in Query Plans
- Do the Lab
- Thank You!

# Pivotal

A NEW PLATFORM FOR A NEW ERA