Greenplum Database Tuning

Pivotal® Greenplum Database

Pivotal

Approaching a Performance Tuning Initiative

The following key points should be followed when tuning:

- Set performance expectations by defining goals
- Set benchmarks
- Know your baseline hardware performance for throughput and capacity
- Know your workload:
  - Heavy usage times
  - Resource contention
  - Data contention
- Focus your optimizations

Pivotal

This slide is intentionally left blank.

Common Causes of Performance Issues

The following are common causes of performance issues:


Hardware issues / failed segments


Resource allocation


Contention between concurrent workloads


Inaccurate database statistics


Uneven data distribution

SQL formulation


Database design

Common Causes of Performance Issues

The following are common causes of performance issues:


Hardware issues / failed segments


Resource allocation


Contention between concurrent workloads


Inaccurate database statistics


Uneven data distribution

SQL formulation


Database design

Common Causes of Performance Issues

The following are common causes of performance issues:


Hardware issues / failed segments


Resource allocation


Contention between concurrent workloads


Inaccurate database statistics


Uneven data distribution

SQL formulation


Database design

Pivotal

Hardware Issues

Common hardware failures include:

- Disk failures
- Host failures
- Network failures
- OS not tuned for Greenplum
- Disk Capacity:
  - 70% maximum recommended
  - VACUUM after updates, deletes and loads
- VACUUM configuration parameters

Pivotal

Hardware Issues – `VACUUM` Configuration Parameters

Set the following `VACUUM` configuration parameters:

- `max_fsm_relations`:
  - This parameter should be set to *tables* + *indexes* + *system tables*
  - Sets the number of relations for which free space will be tracked in the memory free-space map
- `max_fsm_pages`:
  - This parameter is equal to *16 * max_fsm_relations*
  - Sets the number of disk pages for which free space will be tracked

Resource Allocation and Contention

To work around resource allocation issues:

- Greenplum resource queues
  - Limit active queries in the system
  - Limit the size of a query a particular user can run
- Perform admin tasks at low usage times
  - Data loading, ETL
  - `VACUUM` and `ANALYZE`
  - Backups
- Design applications to prevent lock conflicts
  - Concurrent sessions not updating the same data at the same time
- Set resource-related configuration parameters

Setting Resource Related Configuration Parameters

Resource-related configuration parameters include:

- `work_mem = 32MB`
- `maintenance_work_mem = 64MB`
- `shared_buffers = 125MB`

Example: Set and reset a configuration parameter

```
=# SET work_mem TO '200MB';
=# …SQL statements…;
=# RESET work_mem;
```

Example: Set a configuration parameter for a role

```
ALTER ROLE admin SET maintenance_work_mem = 100000;
```

Pivotal

This slide is intentionally left blank.

Setting Memory Management Parameters

Memory management parameters include:

- `statement_mem = 125 MB`
- `max_statement_mem = 2000 MB`
  `(segment_physical_memory/`
  `average_number_concurrent_queries)`
- `gp_vmem_protect_limit = 8192`
  `(X * physical_memory )/primary_segments`

Note: These parameters are used by Greenplum only when `gp_resqueue_memory_policy` is set to `eager_free` or `auto`.

Pivotal

Database Statistics – `ANALYZE`

Greenplum:

- Uses a statistics-based query planner
- Collects information such rows and range of values
- Uses `ANALYZE` to collect statistics. It should be run after:
  - Data loads
  - Restores from backups
  - Changes to schema
  - Inserts, updates, or deletes

Configuring Statistics Collection

Use the following to configure statistics collection:

- `default_statistics_target = 25`
- `gp_analyze_relative_error = .25`
- On specific table columns, run:

```
ALTER TABLE name ALTER column SET STATISTICS #;
```

Greenplum Data Distribution

When working with data:

- Consider your table distribution key
- Check for data skew and avoid, if possible, unbalanced data
- Rebalancing a table if necessary

Greenplum Data Distribution – Consider the Table Distribution Key

When deciding on the table distribution key, look for:

- Even data distribution, where:
  - All segments should contain equal portions of data
  - The distribution key is unique for each record
- Local over distributed operations, where:
  - It is faster if the work can be performed at the segment level
  - A common distribution key improves joining or sorting
  - Local operations can be 5 times faster than distributed operations
- Even query processing, where:
  - All segments handle an equal amount of the query workload
  - Distribution policy and query predicates are well matched

Pivotal

This slide is intentionally left blank.

Greenplum Data Distribution – Check for Data Skew

Check for data skew using:

- `gp_toolkit.gp_skew_coefficients`
- `gp_toolkit.gp_skew_idle_fractions`
- System tools using `gpssh` to run them on multiple systems:
  - `top`
  - `iostat`

Pivotal

Greenplum Data Distribution – Rebalancing a Table

Rebalancing a table can be performed with the following:

- Change the distribution policy to a different column and redistribute the table and child tables:

```
ALTER TABLE sales SET DISTRIBUTED BY (customer_id);
```

- Redistribute table data to correct data skew:

```
ALTER TABLE sales SET WITH (REORGANIZE=TRUE);
```

Pivotal

SQL Formulation – General Considerations

When creating your queries:

- Know your data
- Minimize returned rows
- Avoid unnecessary columns in the result set
- Avoid unnecessary tables
- Avoid sorts of large result sets
- Match data types in predicates

SQL Formulation – Greenplum Specific Considerations

Greenplum-specific guidelines for creating queries include:

- Use common distribution keys:
  - For joins and aggregations
  - So most of the work is performed at the segment level
- Consider the table data distribution policy and query predicates:
  - To have segments handle an equal amount of work
  - To provide the best possible performance

Database Design

When considering the database design:

- Select appropriate data types
- Use a denormalized model
- Consider table partitioning
- Reconsider the use of indexes

Database Design – Selecting Appropriate Data Types

When selecting data types, choose a data type:

- That uses the least possible space
- That best constrains the data:
  - Use character data types for strings
  - Use date or timestamp data types for dates
  - Use numeric data types for numbers
  - Use `TEXT` or `VARCHAR` for character data
- Use identical data types for columns used in cross-table joins

This slide is intentionally left blank.

Database Design – Denormalization

Normalization:

- Is the process of eliminating redundancy and improving data organization
- Is used by online transaction processing (OLTP) databases

Denormalization:

- Is used by online analytical processing (OLAP) databases
- Translates into redundant data
- May facilitate ease of use and performance
- Is used by the star schema, where:
  - Data is stored in a central fact table
  - Dimension tables are denormalized
  - Complexity of queries is reduced
  - ETL processing may be required

Pivotal

Database Design – Table Partitioning

Table partitioning:

- Addresses the problem of supporting very large tables
- Divides large tables into smaller, manageable pieces
- Can improve query performance
- Lets the query planner scan only relevant data
- Should be used to help selectively scan data based on query predicates

Database Design – Indexes

If you are considering indexes, use the following guidelines:

- Use sparingly in Greenplum Database
- Test the query workload without indexes
- Ensure any indexes added are used by the query workload
- Verify that indexes improve query performance
- Indexes can improve performance of OLTP type workloads

Database Design – Index Considerations

When incorporating indexes, use the following guidelines:

- Avoid using indexes on frequently updated columns
- Avoid overlapping indexes
- Use bitmap indexes where applicable instead of B-tree
- Drop indexes for loads
- Consider a clustered index
- Configuring index usage with the following:
  `enable_indexscan = on | off`
- Compressed append-optimized tables may benefit from indexes
- If indexing partitioned tables, index columns should not be the same as partition columns

Pivotal

This slide is intentionally left blank.

Tracking Performance Issues

Performance management steps taken:

- Are often reactive
- Can focus efforts on tuning specific workloads
- Can be caused by:
  - Hardware problems
  - System failures
  - Resource contention
- Can be tracked with:
  - `pg_stat_activity`
  - `pg_locks` or `pg_class`
  - Database logs
  - UNIX system utilities

Pivotal

Tracking Performance Issues – `pg_stat_activity` System Catalog View

The `pg_stat_activity` view:

- Is a system catalog view
- Shows one row per server process

All processes that are not IDLE

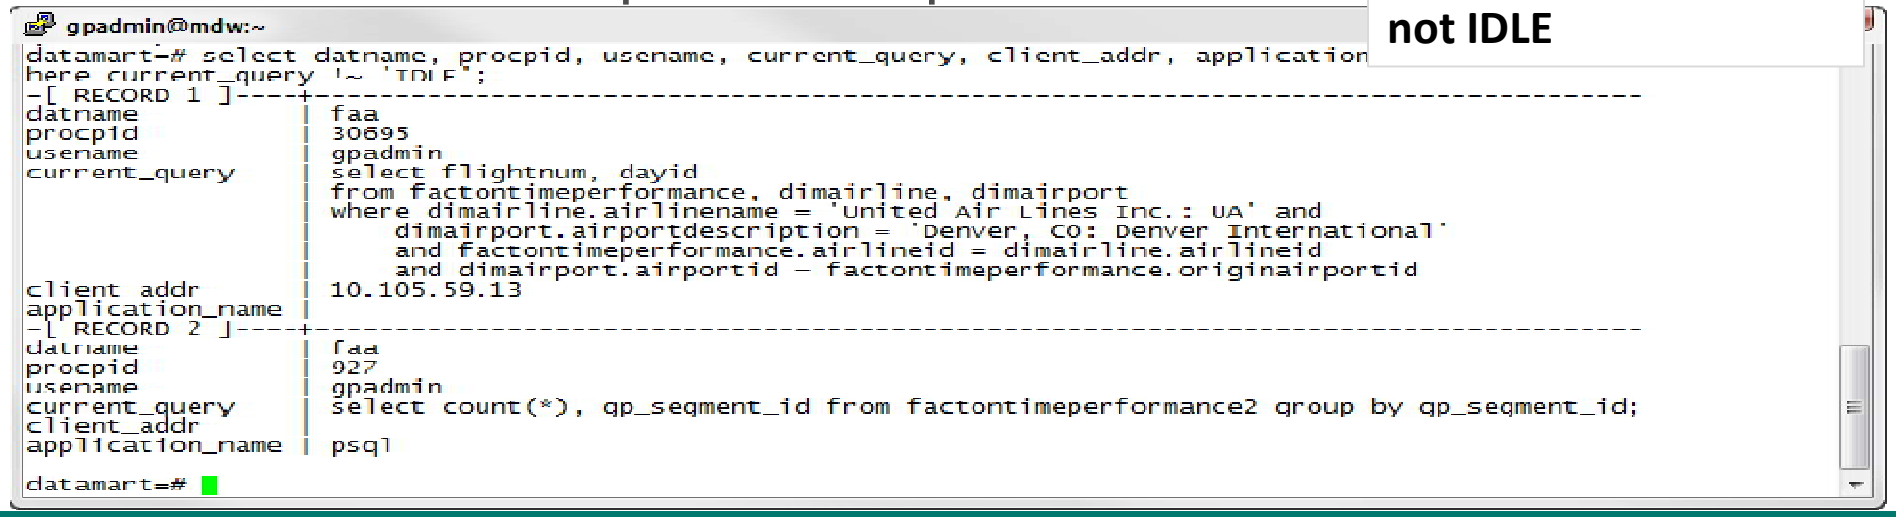

Tracking Performance Issues – `pg_locks` System Activity View

The `pg_locks` view:

- Is a system catalog view
- Lets you view information on outstanding locks
- Can help identify contention between sessions
- Provides a global view of all locks in the database system
- Can be joined to `pg_class.oid` for relations in the current database
- Can have the `pid` column joined to `pg_stat_activity.procid` for more session information

Pivotal

Tracking Performance Issues – `pg_locks` System Activity View (Cont)

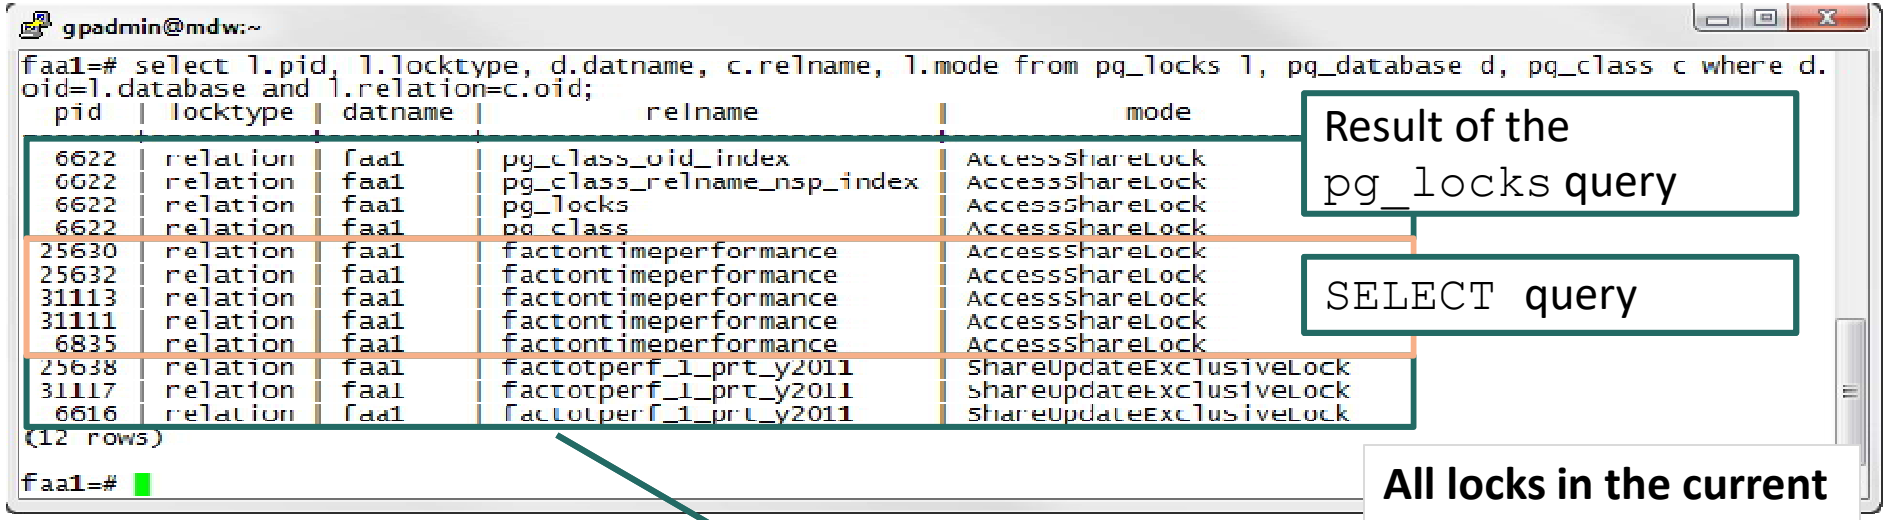

Tracking Performance Issues – Greenplum Database Log Files

Log files:

- Can be found for the master and segments
- Is located in the data directory location of the instance
- Can be accessed with:
  - `gpstate -l` to get the location of log files
  - `gpstate -e` to list the last lines of the log files

Tracking Performance Issues – UNIX System Utilities

System monitoring tools:

- Include:
  - `ps`
  - `top`
  - `iostat`
  - `vmstat`
  - `netstat`
- Help to:
  - Identify processes running on the system
  - Identify the most resource intensive tasks
- Can help identify queries overloading system resources
- Can be run on several hosts at once using `gpssh`

Wrapping Up

In this module we covered:

- Key steps in approaching performance tuning
- Common factors that can affect performance
- Best practices, commands and tools to help tune the Greenplum Database system


Pivotal
A NEW PLATFORM FOR A NEW ERA