

# Pivotal

A NEW PLATFORM FOR A NEW ERA

# Query Profiling



Pivotal® **Greenplum**  
**Database**

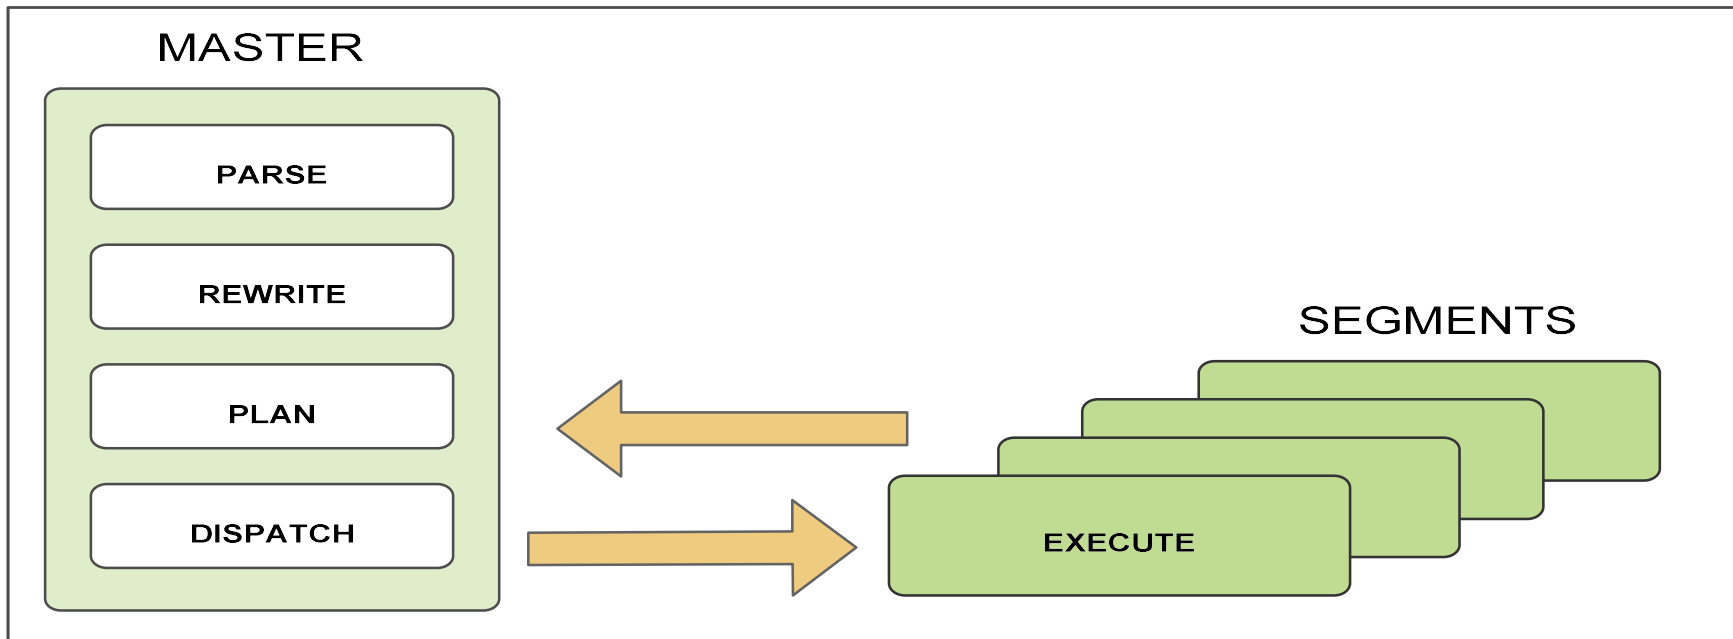
# Agenda

- The Query Process
- Profiling
- Structure of a Query Plan
- Identify Tuning Opportunities
- Work the lab

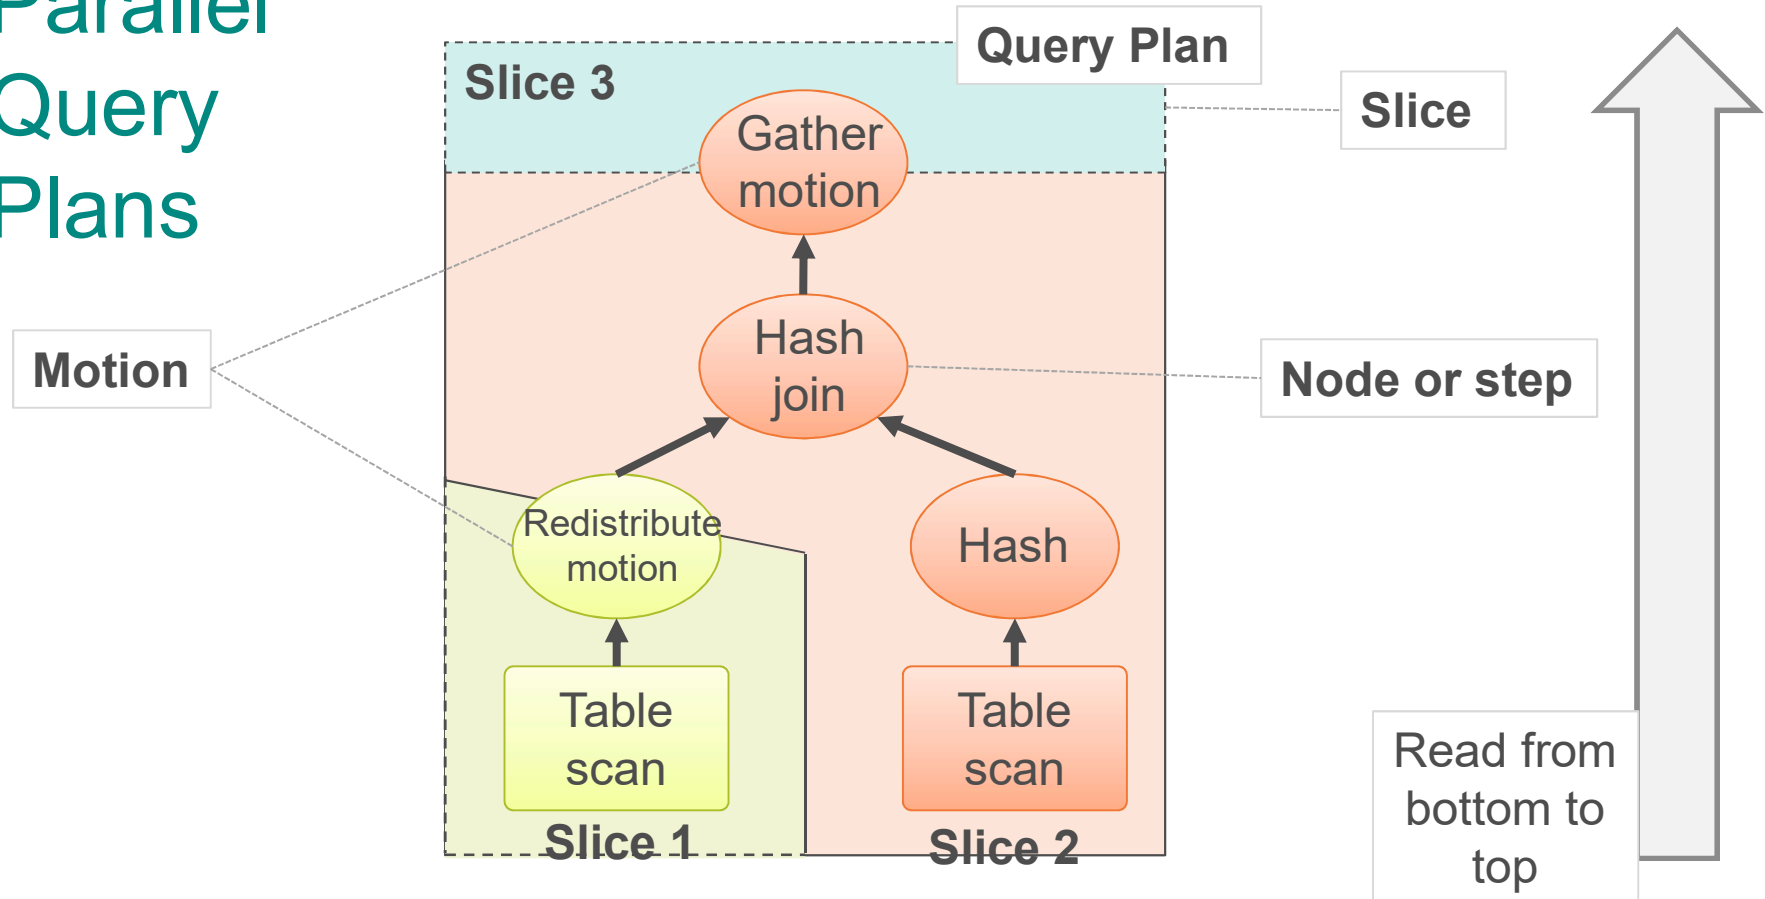
# Query Profiling

- Using query plans to identify tuning opportunities
- Assists you in addressing performance issues
- *Am I getting my results as quickly as possible?*

# The Query Process




# Parallel Query Plans

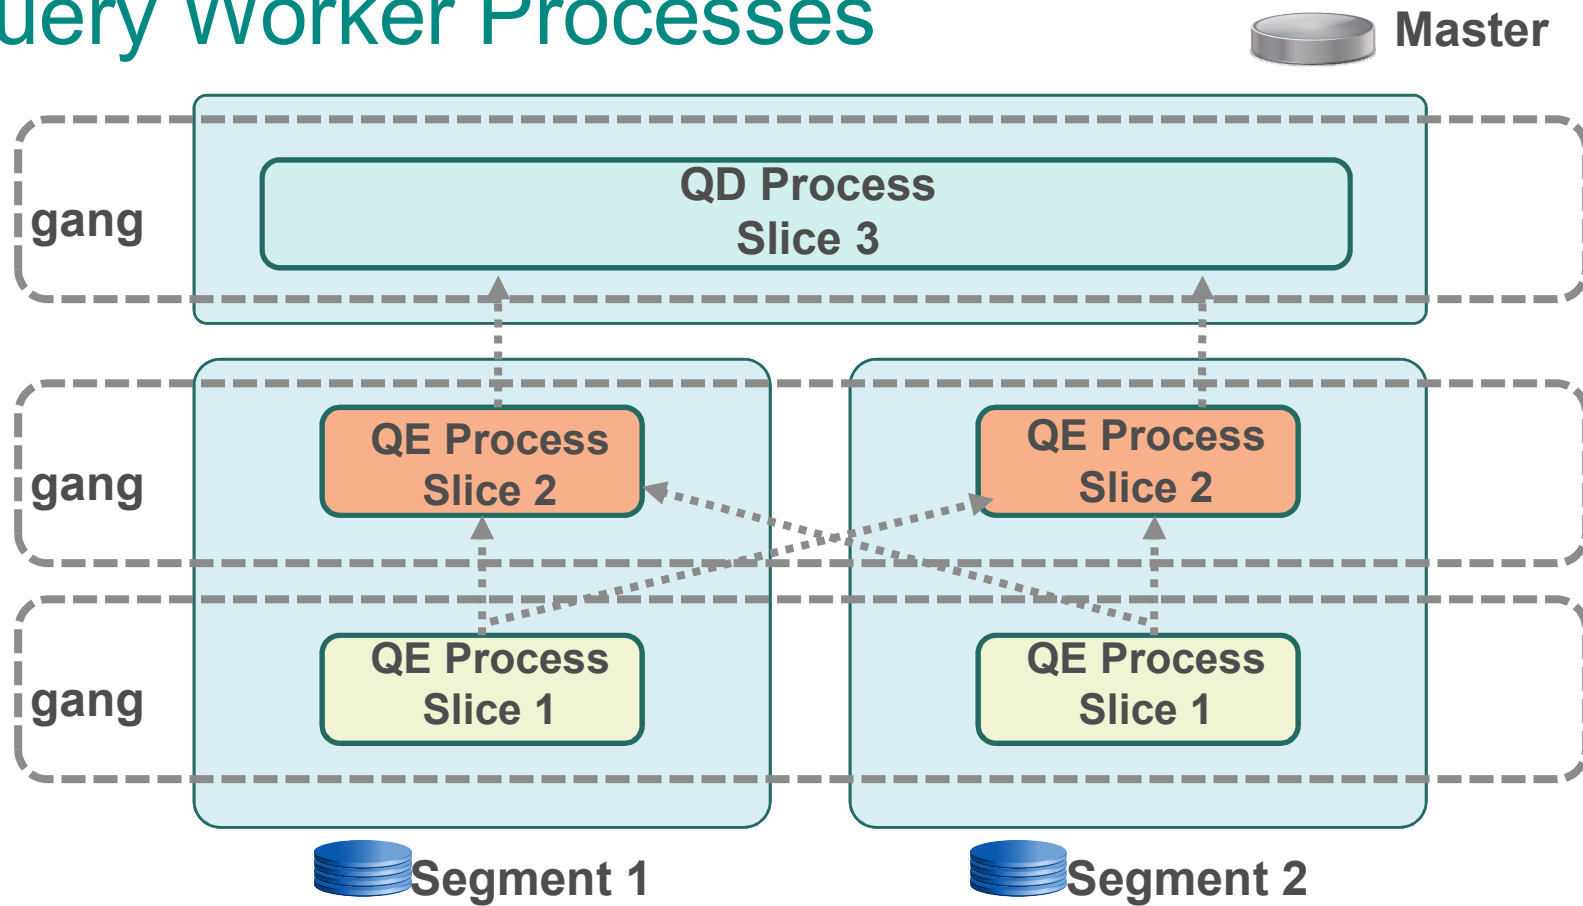


# Parallel Query Plans

```
SELECT customer, amount  
FROM sales JOIN customer  
  USING (cust_id)  
WHERE ship_date=04302016;
```



# Query Worker Processes





# Query Profiling – Addressing Key Issues

- Operations that are taking a long time
- Are the planner's estimates close to reality?
- Is the planner applying selective predicates early?
- Is the planner choosing the best join order?

# Query Profiling – Addressing Key Issues

- Is the planner selectively scanning partitioned tables?
- Is the planner choosing hash aggregate and hash join operations where applicable?
- Is there sufficient work memory?

# Viewing the Query Plan

To see the plan for a query, use:

- `EXPLAIN <query>`
- `EXPLAIN ANALYZE <query>`

Query plans:

- Are read from bottom to top
- Include motions, such as Gather, Redistribute, Broadcast on
  - Joins
  - Sorts
  - Aggregations

The following metrics are given for each operation:

- Cost (units of disk page fetches)
- Rows (rows output by this node)
- Width (average size, in bytes, of a node's output)

# Reading the EXPLAIN Output

SQL query

```
EXPLAIN SELECT * FROM names WHERE name = 'Joelle';
```

----- QUERY PLAN -----

- Gather motion

Cost, rows, and width

```
Gather Motion 2:1 (slice1) (cost=0.00..20.88 rows=1 width=13)  
  -> Seq Scan on 'names' (cost=0.00..20.88 rows=1 width=13)  
    Filter: name::text ~~ 'Joelle'::text
```

Filter condition

Pivotal™

# Reading EXPLAIN ANALYZE Output

```
EXPLAIN ANALYZE SELECT * FROM names WHERE name = 'Joelle';
```

1 row returned  
to master

----- QUERY PLAN -----

```
Gather Motion 2:1 (slice1) (cost=0.00..20.88 rows=1 width=13)
```

```
  recv: Total 1 rows with 0.305 ms to first row, 0.537 ms to  
  end.
```

```
    -> Seq Scan on 'names' (cost=0.00..20.88 rows=1
```

```
      Total 1 rows (seg0) with 0.255 ms to s
```

```
to end.
```

1 segment  
returned rows

```
      Filter: name::text ~~ 'Joelle'::text
```

```
22.548 ms elapsed
```

Actual time to  
run the query

Pivotal™

# Review

- The Query Process
- Profiling
- Structure of a Query Plan
- Identify Tuning Opportunities
- Work the lab

# Pivotal

A NEW PLATFORM FOR A NEW ERA